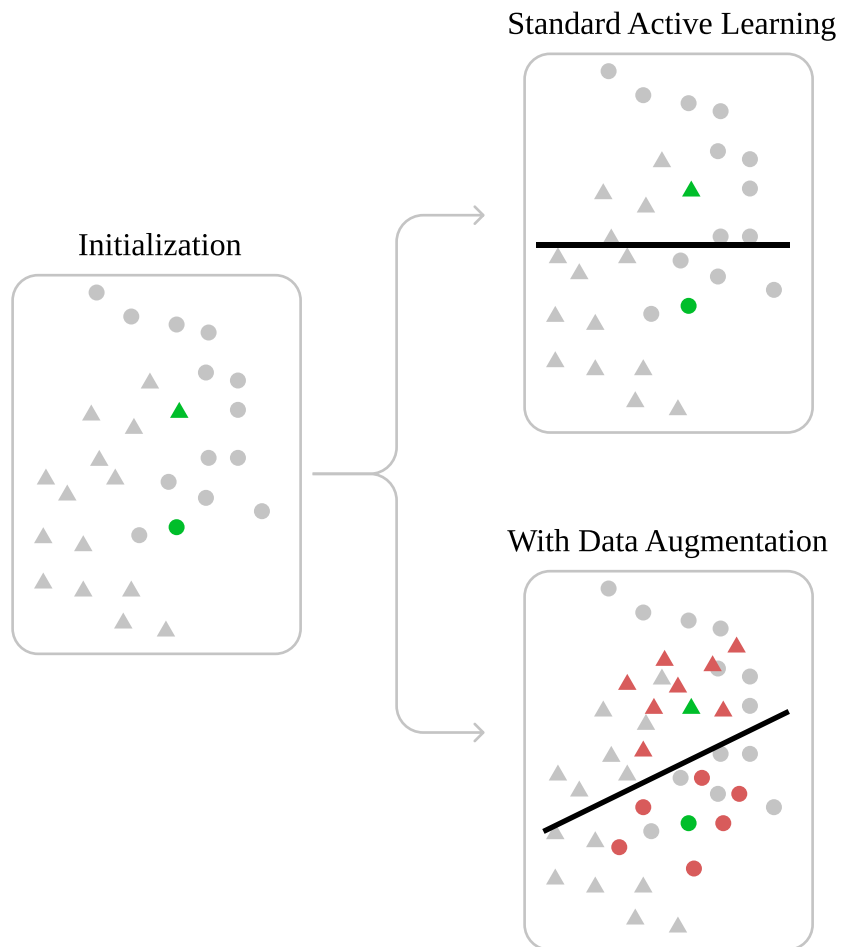


Graphical Abstract

Improving Active Learning Performance Through the Use of Data Augmentation

Joao Fonseca, Fernando Bacao



Highlights

Improving Active Learning Performance Through the Use of Data Augmentation

Joao Fonseca, Fernando Bacao

- We propose a new Active Learning framework that leverages hyperparameter optimization and data augmentation techniques;
- The use of data augmentation in Active Learning is sufficient to substantially improve the performance of an Active Learner, regardless of the choice of dataset/domain, classifier or metric.
- In most scenarios, the proposed method outperformed classifiers trained in fully supervised settings while using less data.

Improving Active Learning Performance Through the Use of Data Augmentation

Joao Fonseca^a, Fernando Bacao^a

^a*NOVA Information Management School, Universidade Nova de Lisboa, Campus de Campolide, Lisboa, 1070-312, Lisboa, Portugal*

Abstract

Active Learning (AL) is a technique that is used to iteratively select unlabeled observations out of a large pool of unlabeled data to be labeled by a supervisor. Its focus is to find the unlabeled observations that, once labeled, will maximize the informativeness of the training dataset. However, the manual labeling of observations involves human resources with domain expertise, making it an expensive and time-consuming task. The literature describes various methods to improve the effectiveness of this process, but there is little research developed around the usage of artificial data sources in AL. In this paper we propose a new framework for AL, which allows for an effective use of artificial data. Our method implements a data augmentation policy that optimizes the generation of artificial instances to improve the AL process. We compare the proposed method to the standard framework as well as another active learning method that uses data augmentation. The models' performance was tested using 4 different classifiers, 2 AL-specific performance metrics and 3 classification performance metrics over 10 different datasets. We show that the proposed framework, using data augmentation, significantly improves the performance of AL, both in terms of classification performance and data selection efficiency.

Keywords: Active Learning, Data Augmentation, Oversampling

1. Introduction

The importance of training robust ML models with minimal data requirements is substantially increasing [1, 2, 3]. Although the growing amount of

valuable data sources and formats being developed and explored is affecting various domains [4], this data is often unlabeled. Only a small amount of the data being produced and stored can be useful for supervised learning tasks. Additionally, it’s important to note that labeling data for specific Machine Learning (ML) projects is often difficult and expensive, especially when data-intensive ML techniques are involved (*e.g.*, Deep Learning classifiers) [1]. In this scenario, labeling the full dataset becomes impractical, time-consuming, and expensive. Two different ML techniques attempt to address this problem: Semi-Supervised Learning (SSL) and Active Learning (AL). Even though they address the same problem, the two follow different approaches. SSL focuses on observations with the most certain predictions, whereas AL focuses on observations with the least certain predictions [5].

SSL attempts to use a small, predefined set of labeled and unlabeled data to produce a classifier with superior performance. This method uses the unlabeled observations to help define the classifier’s decision boundaries [6]. Simultaneously, the amount of labeled data required to reach a given performance threshold is also reduced. It is a special case of ML because it falls between the supervised and unsupervised learning perspectives. AL, instead of optimizing the informativeness of an existing training set, it expands the dataset to include the most informative and/or representative observations [7]. It is an iterative process where a supervised model is trained and simultaneously identifies the most informative unlabeled observations to increase the performance of that classifier. The combination of SSL with AL has been explored in the past, achieving state-of-the-art results [8].

Several studies have pointed out the limitations of AL within an Imbalanced Learning context [9]. With imbalanced data, AL approaches frequently have low performance, high computational time, or data annotation costs. Studies addressing this issue tend to adopt classifier-level modifications, such as the Weighted Extreme Learning Machine [9, 10, 11]. However, classifier or query function-level modifications (See Section 2.1) have limited applicability since a universally good AL strategy has not been found [7]. Other methods address imbalance learning by weighing the observations as a function of the observation’s class imbalance ratio [12]. Alternatively, other methods reduce the imbalanced learning bias by combining Informative and Representative-based query approaches (see Section 2.1) [13]. Another approach to deal with imbalanced data and data scarcity, in general, is data augmentation. This approach has the advantage of being classifier-agnostic, potentially reduces the imbalanced learning bias, and also works as a regularization method in

data-scarce environments, such as AL implementations [14]. However, most recent studies improve the AL performance by modifying the design/choice of the classifier and query functions used.

The usage of data augmentation in AL is not new. The literature found on the topic (see Section 2.3) focuses on either image classification or Natural Language Processing and uses Deep Learning-based data augmentation to improve the performance of neural network architectures in AL. These methods, although showing promising results, represent a limited perspective of the potential of data augmentation in a real-world setting. First, using Deep Learning in an iterative setting requires access to significant computational power. Second, these models tend to use sophisticated data augmentation methods, whose implementation may not be accessible to the non-sophisticated user. Third, the studies found on the topic are specific to the domain, classifier and data augmentation method. Consequently, the direct effect of data augmentation is unclear: these studies implement different neural network-based techniques for different classification problems, whose performance may be attributed to various elements within the AL framework.

In this study, we explore the effect of data augmentation in AL in a context-agnostic setting, along with two different data augmentation policies: oversampling (where the amount of data generated for each class equals the amount of data belonging to the majority class) and non-constant data augmentation policies (where the amount of data generated exceeds the amount of data belonging to the majority class in varying quantities) between iterations. We start by conceptualizing the AL framework and each of its elements, as well as the modifications involved to implement data augmentation in the AL iterative process. We argue that simple, non-domain specific data augmentation heuristics are sufficient to improve the performance of AL implementations, without the need to resort to deep learning-based data augmentation algorithms.

When compared to the standard AL framework, the proposed framework contains two additional components: the Generator and the Hyperparameter Optimizer. We implement a modified version of Geometric Synthetic Minority Oversampling Technique (G-SMOTE) [15] as a data augmentation method with an optimized generation policy (explained in Section 2.2). We also propose a hyperparameter optimization module, which is used to find the best data augmentation policy at each iteration. We test the effectiveness of the proposed method in 10 datasets of different domains. We implement

3 AL frameworks (standard, oversampling and varying data augmentation) using 4 different classifiers, 3 different performance metrics and calculate 2 AL-specific performance metrics.

The rest of this manuscript is structured as follows: Section 2.1 describes the state-of-the-art in AL. Section 2.2 describes the state-of-the-art in Data Augmentation. Section 3 describes the proposed method. Section 4 describes the methodology of the study’s experiment. Section 5 presents the results obtained from the experiment, as well as a discussion of these results. Section 6 presents the conclusions drawn from this study.

2. Background

2.1. Active Learning

This paper focuses on pool-based AL methods as defined in [16]. The goal of AL models is to maximize the performance of a classifier, f_c , while annotating as least observations, x_i , as possible. They use a data pool, \mathcal{D} , where $\mathcal{D} = \mathcal{D}_{lab} \cup \mathcal{D}_{pool}$ and $|\mathcal{D}_{pool}| \gg |\mathcal{D}_{lab}|$. \mathcal{D}_{pool} and \mathcal{D}_{lab} refer to the sets of unlabeled and labeled data, respectively. Having a budget of T iterations (where $t = 1, 2, \dots, T$) and n annotations per iteration, at iteration t , f_c is trained using \mathcal{D}_{lab}^t to produce, for each $x_i \in \mathcal{D}_{pool}^t$, an uncertainty score using an acquisition function $f_{acq}(x_i; f_c)$. These uncertainty scores are used to annotate the n observations with highest ucertainty from \mathcal{D}_{pool}^t to form \mathcal{D}_{new}^t . The iteration ends with the update of $\mathcal{D}_{lab}^{t+1} = \mathcal{D}_{lab}^t \cup \mathcal{D}_{new}^t$ and $\mathcal{D}_{pool}^{t+1} = \mathcal{D}_{pool}^t \setminus \mathcal{D}_{new}^t$ [17, 2]. This process is shown in Figure 1. Before the start of the iterative process, assuming $\mathcal{D}_{lab}^{t=0} = \emptyset$, the data used to populate $\mathcal{D}_{lab}^{t=1}$ is typically collected randomly from $\mathcal{D} = \mathcal{D}_{pool}^{t=0}$ and is labeled by a supervisor [18, 19, 20].

Research focused on AL has typically been focused on the specification of f_{acq} and domain-specific applications. Acquisition functions can be divided into two different categories [21, 22]:

1. Informative-based. These strategies use the classifier’s output to assess the importance of each observation towards the performance of the classifier [23].
2. Representative-based. These strategies estimate the optimal set of observations that will optimize the classifier’s performance [22].

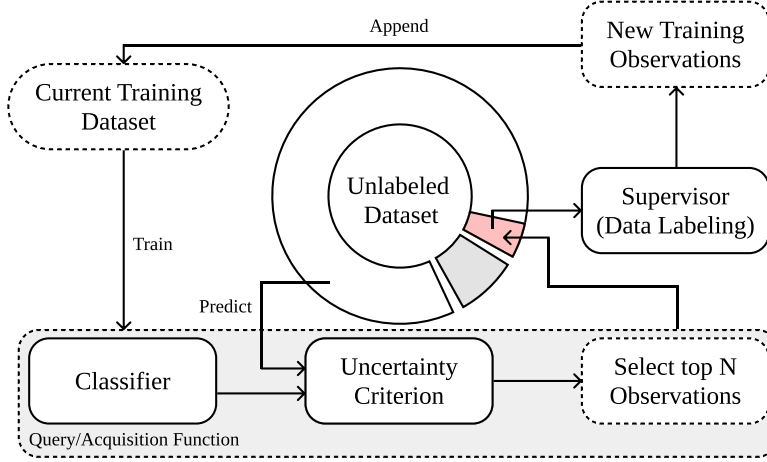


Figure 1: Diagram depicting a typical AL iteration. In the first iteration, the training set collected during the initialization process becomes the “Current Training Dataset”.

114 Although there are significant contributions towards the development of
 115 more robust query functions and classifiers in AL, modifications to AL’s
 116 basic structure is rarely explored. In [19] the authors introduce a loss predic-
 117 tion module in the AL framework to replace the uncertainty criterion. This
 118 model implements a second classifier to predict the expected loss of the un-
 119 labeled observations (using the actual losses collected during the training of
 120 the original classifier) and return the unlabeled observations with the high-
 121 est expected loss. However, this contribution is specific to neural networks
 122 (and more specifically, to deep neural networks) and was only tested for image
 123 classification.

124 2.2. Data Augmentation

125
 126 Data Augmentation methods expand the training dataset by introducing
 127 new and informative observations [24]. The production of artificial data may
 128 be done via the introduction of perturbations on the input [25], feature [26]
 129 or output space [24]. Data Augmentation methods may be divided into two
 130 categories [27]:

- 131 1. Heuristic approaches attempt to generate new and relevant observa-
 132 tions through the application a predefined procedure, usually incorpo-
 133 rating some degree of randomness [28]. Since these methods typically

occur in the input space, they require less data and computational power when compared to Neural Network methods.

2. Neural Network approaches, on the other hand, map the original input space into a lower-dimensional representation, known as the feature space [26]. The generation of artificial data occurs in the feature space and is reconstructed into the input space. Although these methods allow the generation of less noisy data in high-dimensional contexts and more plausible artificial data, they are significantly more computationally intensive.

While some techniques may depend on the domain, others are domain-agnostic. For example, Random Erasing [25], Translation, Cropping and Flipping are examples of image data-specific augmentation methods. Other methods, such as autoencoders, may be considered domain agnostic.

2.3. Data Augmentation in Active Learning

The standard AL model can be complemented with a data augmentation function, $f_{aug}(x_i; \tau)$, where τ defines the augmentation policy. In this context, τ refers to the transformation applied and its hyperparameters and $f_{aug}(x; \tau) : \mathcal{D} \rightarrow \mathcal{D}_{aug}(\mathcal{D})$ produces a modified observation, $\tilde{x} \in \mathcal{D}_{aug}(\mathcal{D})$ where $\mathcal{D}_{aug}(\mathcal{D})$ is the set of modified observations. This involves the usage of a new set of data, $\mathcal{D}_{train}^t = \mathcal{D}_{lab}^t \cup \mathcal{D}_{aug}^t$, to train the classifier.

As found in Section 2.1, improvements proposed in the AL framework are mostly focused on modifications of the classifier or query strategy. Furthermore, the few recent AL contributions implementing data augmentation were all (except one) applied to the computer vision or natural language processing (NLP) realm.

The only AL model found that uses data augmentation outside of the computer vision or NLP domains uses a pipelined approach, described in [18]. In this study, the AL model proposed is applied for tabular data using an oversampling data augmentation policy (*i.e.*, the artificial data was generated only to balance the target class frequencies).

the AL model proposed was applied in a Land Use/Land Cover context with specific characteristics that are not necessarily found in other supervised learning problems. Specifically, these types of datasets are high dimensional and have limited data variability within each class (*i.e.*, cohesive spectral

signatures within classes) due to their geographical proximity. Furthermore, this method does not allow augmentation policy optimization (i.e., every hyperparameter has to be hard coded a priori).

The Bayesian Generative Active Deep Learning (BGDAL) [29] is another example of a pipelined combination of f_{acq} and f_{aug} , applied image classification. BGDAL uses a Variational AutoEncoder (VAE) architecture to generate artificial observations. However, the proposed model is computationally expensive, requires a large data pool to train the VAE, and is not only dependent on the quality of the augmentations performed, but also on the performance of the discriminator and classifiers used.

The method proposed in [14], Look-Ahead Data Acquisition for Deep Active Learning, implement data augmentation to train a deep learning classifier. However, adapting existing AL applications to use this approach is often impractical and implies the usage of image data, since the augmentations used are image data specific and occur on the unlabeled observations, before the unlabeled data selection.

The Variational Adversarial Active Learning (VAAL) model [30] is a deep AL approach to image classification that uses as inputs the embeddings produced by a VAE into a secondary classifier, working as f_{acq} , to predict if $x_i \in \mathcal{D}$ belongs to \mathcal{D}_{pool} . The n true positives with the highest certainty are labeled by the supervisor and \mathcal{D}_{pool} and \mathcal{D}_{lab} are updated as described in Section 2.1. The Task-aware VAAL model [31] extends the VAAL model by introducing a ranker, which consists of the Learning Loss module introduced in [19]. These models use data augmentation techniques to train the different neural network-based components of the proposed models. However, the AL components used are specific image classification, computationally expensive and the analysis of the effect of data augmentation in these AL models is not discussed.

In [32] the proposed AL method was designed specifically for image data classification, where a deep learning model was implemented as a classifier, but its architecture is not described, the augmentation policies used are unknown and the results reported correspond to single runs of the discussed model. The remaining AL models found implementing data augmentation were introduced for NLP applications, in [33, 34]. However, these methods were designed for specific applications within that domain and are not necessarily transferable to other domains or tasks.

205 3. Proposed Method

206
207 Based on the literature found on AL, most of the contributions and
208 novel implementations of AL algorithms focused on the improvement of the
209 choice/architecture of the classifier or the improvement of the uncertainty
210 criterion. In addition, the resulting classification performance of AL-trained
211 classifiers is frequently inconsistent and marginally improve the classifica-
212 tion performance when compared to classifiers trained over the full training
213 set. In addition, there is also a significant variability of the data selection
214 efficiency during different runs of the AL iterative process [18].

215 This paper provides a context-agnostic AL framework towards the inte-
216 gration of Data Augmentation within AL, with the following contributions:

- 217 1. Improvement of the AL framework by introducing a parameter tuning
218 stage using only the labeled dataset available at the current iteration
219 (*i.e.*, no labeled hold-out set is needed).
- 220 2. Generalization of the generator module proposed in [18] from oversam-
221 pling techniques to any other data augmentation mechanism and/or
222 policy.
- 223 3. Implementation of data augmentation outside of the Deep AL realm,
224 which was not previously found in the literature.
- 225 4. Analysis of the impact of Data Augmentation and Oversampling in AL
226 over 10 different datasets of different domains, while comparing them
227 with the standard AL framework.

228 The proposed AL framework is depicted in Figure 2. The generator el-
229 ement becomes an additional source of data and is expected to introduce
230 additional data variability into the training dataset. This should allow the
231 classifier to generalize better and perform more consistently over unseen ob-
232 servations. However, in this scenario, the amount of data to generate per
233 class at each iteration is unknown. Consequently, the hyperparameter tun-
234 ing step was introduced to estimate the optimal data augmentation policy
235 at each iteration. In our implementation, this step uses the current training
236 dataset to perform an exhaustive search over specified parameters of the gen-
237 erator, tested over a 5-fold cross validation method. The best augmentation

238 policy found is used to train the iteration’s classifier in the following step.
 239 This procedure is described in Algorithm 1.

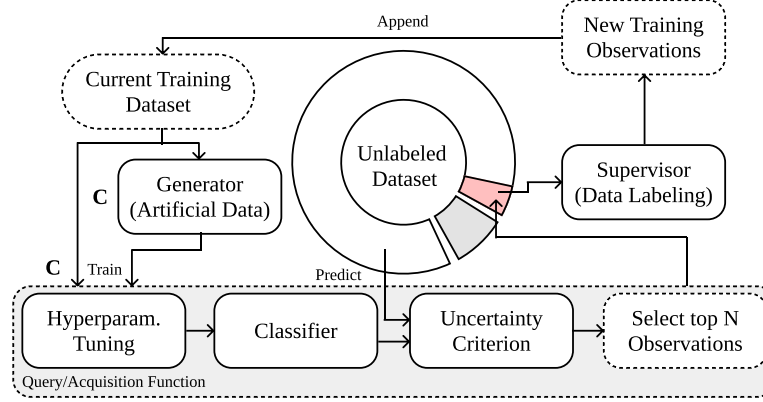


Figure 2: Diagram depicting the proposed AL iteration. The proposed modifications are marked with a boldface “C”.

240 To show the effectiveness of data augmentation in an AL implementation,
 241 we implemented a simple modification in the selection mechanism of the G-
 242 SMOTE algorithm. We use the uncertainties produced by f_{acq} to compute
 243 the probabilities of observations to be selected for augmentation, as an ad-
 244 ditional parameter. This modification is described in Algorithm 2

245 This modification facilitates the usage of G-SMOTE beyond its original
 246 oversampling purposes. In this paper, the data augmentation strategies used
 247 ensure the all the class frequencies are balanced. Furthermore, the amount of
 248 artificial data produced for each class is defined by the *augmentation factor*,
 249 α_{af} , which represents a percentage of the majority class C_{maj} (e.g., an aug-
 250 mentation factor of 1.2 will ensure there are $count(C_{maj}) \times 1.2$ observations
 251 in every class). In this paper’s experiment, the data generation mechanism
 252 is similar to the one in [18]. This allows the direct comparison of the two
 253 frameworks and establish a causality of the performance variations to the
 254 data generation mechanism (i.e., augmentation vs normal oversampling) and
 255 hyperparameter tuning steps.

Algorithm 1: Proposed AL Framework (Single iteration)

Given: $t \geq 1$, performance metric f_{pm}
Input: $\mathcal{D}_{pool}, \mathcal{D}_{lab}, f_c, f_{aug}, f_{acq}, \tau_{grid}, k, n$
Output: $\mathcal{D}_{pool}, \mathcal{D}_{lab}$

```

1 Function ParameterTuning( $f_c, f_{aug}, \tau_{grid}, \mathcal{D}_{lab}, k$ ):
2    $p \leftarrow 0$ 
3    $\tau \leftarrow \emptyset$ 
4    $\{\mathcal{D}_{lab}^1, \dots, \mathcal{D}_{lab}^k\} \leftarrow \mathcal{D}_{lab}$       //  $\mathcal{D}_{lab}^n \cap \mathcal{D}_{lab}^m = \emptyset, \forall (n, m) \in 1, \dots, k$ 
5   forall  $\tau' \in \tau_{grid}$  do
6      $p' \leftarrow \emptyset$ 
7     forall  $\mathcal{D}_{lab}^i \in \{\mathcal{D}_{lab}^1, \dots, \mathcal{D}_{lab}^k\}$  do
8        $\mathcal{D}'_{test} \leftarrow \mathcal{D}_{lab}^i$ 
9        $\mathcal{D}'_{train} \leftarrow \mathcal{D}_{lab} \setminus \mathcal{D}_{lab}^i$ 
10       $\mathcal{D}'_{train} \leftarrow f_{aug}(\mathcal{D}'_{train}; \tau')$ 
11      train  $f_c$  using  $\mathcal{D}'_{train}$ 
12       $p' \leftarrow p' \cup \{f_{pm}(f_c(\mathcal{D}_{test}))\}$ 
13     $p' \leftarrow \frac{\sum_{x_i \in p'} x_i}{k}$ 
14    if  $p' > p$  then
15       $p \leftarrow p'$ 
16       $\tau \leftarrow \tau'$ 
17  return  $\tau$ 
18 begin
19   $\tau \leftarrow \text{ParameterTuning}(f_c, f_{aug}, \tau_{grid}, \mathcal{D}_{lab}, k)$ 
20   $\mathcal{D}_{train} \leftarrow f_{aug}(\mathcal{D}_{lab}; \tau)$ 
21  train  $f_c$  using  $\mathcal{D}_{train}$ 
22   $\mathcal{D}_{new} = \arg \max_{\mathcal{D}'_{pool} \subset \mathcal{D}_{pool}, |\mathcal{D}'_{pool}|=n} \sum_{x \in \mathcal{D}'_{pool}} f_{acq}(x; f_c)$ 
23  annotate  $\mathcal{D}_{new}$ 
24   $\mathcal{D}_{pool} \leftarrow \mathcal{D}_{pool} \setminus \mathcal{D}_{new}$ 
25   $\mathcal{D}_{lab} \leftarrow \mathcal{D}_{lab} \cup \mathcal{D}_{new}$ 

```

Algorithm 2: G-SMOTE Modified for Data Augmentation in AL

Given: $t \geq 1$, $\mathcal{D}_{lab}^t \neq \emptyset$, $\mathcal{D}_{lab} = \mathcal{D}_{lab}^{min} \cup \mathcal{D}_{lab}^{maj}$, $GSMOTE$
Input: \mathcal{D}_{pool}^t , \mathcal{D}_{lab}^t , f_c^{t-1} , f_{acq} , τ
Output: \mathcal{D}_{train}^t

```

1 Function  $DataSelection(\mathcal{D}_{lab}^t, f_{acq}, f_c^{t-1})$ :
2    $U \leftarrow \emptyset$ 
3    $P \leftarrow \emptyset$ 
4    $p_s \sim \mathcal{U}(0, 1)$ 
5   forall  $x_i \in \mathcal{D}_{lab}^t$  do
6      $u_{x_i} \leftarrow f_{acq}(x_i; f_c^{t-1})$ 
7      $U \leftarrow U \cup \{u_{x_i}\}$ 
8   forall  $u_{x_i} \in U$  do
9      $p_{x_i} \leftarrow \frac{u_{x_i}}{\sum U} + \sum P$ 
10     $P \leftarrow P \cup \{p_{x_i}\}$ 
11     $i \leftarrow \operatorname{argmax}(P < p_s)$ 
12    return  $i$ -th element in  $\mathcal{D}_{lab}^t$ 
13 begin
14    $\mathcal{D}_{aug}^{min} \leftarrow \emptyset$ 
15    $\mathcal{D}_{aug}^{maj} \leftarrow \emptyset$ 
16    $\alpha_{af}, \alpha_{trunc}, \alpha_{def} \leftarrow \tau$ 
17    $N \leftarrow \operatorname{count}(C_{maj}) \times \alpha_{af}$ 
18   forall  $\mathcal{D}'_{aug} \in \{\mathcal{D}_{aug}^{min}, \mathcal{D}_{aug}^{maj}\}$ ,  $\mathcal{D}'_{lab} \in \{\mathcal{D}_{lab}^{min}, \mathcal{D}_{lab}^{maj}\}$  do
19     while  $|\mathcal{D}'_{aug}| < N$  do
20        $x_{center} \leftarrow DataSelection(\mathcal{D}'_{lab}, f_{acq}, f_c^{t-1})$ 
21        $x_{gen} \leftarrow GSMOTE(x_{center}, \mathcal{D}_{lab}^t, \alpha_{trunc}, \alpha_{def})$ 
22        $\mathcal{D}'_{aug} \leftarrow \mathcal{D}'_{aug} \cup \{x_{gen}\}$ 
23    $\mathcal{D}_{aug} \leftarrow \mathcal{D}_{aug}^{min} \cup \mathcal{D}_{aug}^{maj}$ 
24    $\mathcal{D}_{train}^t \leftarrow \mathcal{D}_{lab}^t \cup \mathcal{D}_{aug}$ 

```




Figure 3: Simplified diagrams highlighting the differences between the proposed and standard AL iterations. The proposed modifications are highlight in red and marked with a boldface “C”.

256 The comparison of diagrams between the proposed and standard AL
 257 frameworks is shown in Figure 3. In the proposed framework, we (1) gen-
 258 eralize the generator module to accept any data augmentation method or
 259 policy and (2) a hyperparameter tuning module to estimate the optimal
 260 data augmentation policy. This framework was designed to be task-agnostic.
 261 Specifically, any data augmentation method (domain specific or not) may be
 262 used, as well as any other parameter search method. It is also expected to
 263 be compatible with other AL modifications, including the ones that do not
 264 affect solely the classifier or uncertainty criterion, such as the one proposed
 265 in [19].

266 4. Methodology

267
 268 This section describes the different elements included in the experimental
 269 procedure. The datasets used were acquired in open data repositories and its
 270 sources and preprocessing steps are defined in Subsection 4.1. The choice of
 271 classifiers used in the experiment are defined in Subsection 4.2. The metrics
 272 chosen to measure AL performance and overall classification performance
 273 are defined in Subsection 4.3. The experimental procedure is described in
 274 Subsection 4.4. The implementation of the experiment and resources used
 275 to do so are described in Subsection 4.5.

276 The methodology developed serves 2 purposes: (1) Compare classification
 277 performance once all the AL procedures are completed (*i.e.*, optimal perfor-
 278 mance of a classifier trained via iterative data selection) and (2) Compare the

amount of data required to reach specific performance thresholds (*i.e.*, number of AL iterations required to reach similar classification performances).

4.1. Datasets

The datasets used to test the proposed method are publicly available in open data repositories. Specifically, they were retrieved from OpenML and the UCI Machine Learning Repository. They were chosen considering different domains of application, imbalance ratios, dimensionality and number of target classes, all of them focused on classification tasks. The goal is to demonstrate the performance of the different AL frameworks in various scenarios and domains. The data preprocessing approach was similar across all datasets. Table 1 describes the key properties of the 10 preprocessed datasets where the experimental procedure was applied.

Dataset	Features	Instances	Minority instances	Majority instances	IR	Classes
Image Segmentation	14	1155	165	165	1.0	7
Mfeat Zernike	47	1994	198	200	1.01	10
Texture	40	1824	165	166	1.01	11
Waveform	40	1666	551	564	1.02	3
Pendigits	16	1832	176	191	1.09	10
Vehicle	18	846	199	218	1.1	4
Mice Protein	69	1073	105	150	1.43	8
Gas Drift	128	1987	234	430	1.84	6
Japanese Vowels	12	1992	156	323	2.07	9
Baseball	15	1320	57	1196	20.98	3

Table 1: Description of the datasets collected after data preprocessing. The sampling strategy is similar across datasets. Legend: (IR) Imbalance Ratio

The data preprocessing pipeline is depicted as a flowchart in Figure 4. The missing values are removed from each dataset by removing the corresponding observations. This ensures that the input data in the experiment is kept as close to its original form as possible. The non-metric features (*i.e.*, binary, categorical and ordinal variables) were removed since the application of G-SMOTE is limited to continuous and discrete features. The datasets containing over 2000 observations were downsampled in order to maintain the datasets to a manageable size. The data sampling procedure preserves

the relative class frequency of the dataset, in order to maintain the Imbalance Ratio (IR) originally found in each dataset (where $IR = \frac{\text{count}(C_{maj})}{\text{count}(C_{min})}$). The remaining features of each dataset are scaled to the range of $[-1, 1]$ to ensure a common range across features.

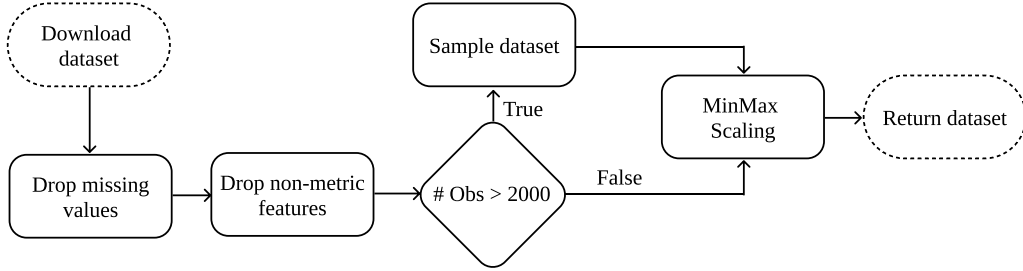


Figure 4: Data preprocessing pipeline.

The preprocessed datasets were stored into a SQLite database file and is available along with the experiment’s source code in the GitHub repository of the project (see Subsection 4.5).

4.2. Machine Learning Algorithms

We used a total of 4 classification algorithms and a heuristic data augmentation mechanism. The choice of classifiers was based on the popularity and family of the classifiers (tree-based, nearest neighbors-based, ensemble-based and linear models). Our proposed method was tested using a Decision Tree (DT) [35], a K-nearest neighbors classifier (KNN) [36], a Random Forest Classifier (RF) [37] and a Logistic Regression (LR) [38]. Since the target variables are multi-class, the LR classifier was implemented using the one-versus-all approach. The predicted class is assigned to the label with the highest likelihood.

The oversampler G-SMOTE was used as a data augmentation method. The typical data generation policy of oversampling methods is to generate artificial observations on non-majority classes such that the number of majority class observations matches those of each non-majority class. We modified this data generation policy to generate observations for all classes, as a percentage of the number of observations in the majority class. In addition, the

original G-SMOTE algorithm was modified to accept data selection probabilities based on classification uncertainty. These modifications are discussed in Section 3.

Every AL procedure was tested with different selection criteria: Random Selection, Entropy and Breaking Ties. The baseline used is the standard AL procedure. As a benchmark, we add the AL procedure using G-SMOTE as a normal oversampling method, as proposed in [18]. Our proposed method was implemented using G-SMOTE as a data augmentation method to generate artificial observations for all classes, while still balancing the class distribution, as described in Section 3.

4.3. Evaluation Metrics

Considering the imbalanced nature of the datasets used in the experiment, commonly used performance metrics such as Overall Accuracy (OA), although being intuitive to interpret, are insufficient to quantify a model’s classification performance [39]. The Cohen’s Kappa performance metric, similar to OA, is also biased towards high frequency classes since its definition is closely related to the OA metric, making its behavior consistent with OA [40]. However, these metrics remain popular choices for the evaluation of classification performance. Other performance metrics like $Precision = \frac{TP}{TP+TN}$, $Recall = \frac{TP}{TP+FN}$ or $Specificity = \frac{TN}{TN+FP}$ are calculated as a function of True/False Positives (TP and FP) and True/False Negatives (TN and FN) and can be used at a per-class basis instead. In a multiple dataset scenario with varying amount of target classes and meanings, comparing the performance of different models using these metrics becomes impractical.

Based on the recommendations found in [39, 41], we used 2 metrics found to be less sensitive to the class imbalance bias, along with OA as a reference for easier interpretability:

- The Geometric-mean scorer (G-mean) consists of the geometric mean of Specificity and Recall [41]. Both metrics are calculated in a multiclass context considering a one-versus-all approach. For multiclass problems, the G-mean scorer is calculated as its average per class values:

$$G\text{-mean} = \sqrt{Sensitivity \times Specificity}$$

- The F-score metric consists of the harmonic mean of Precision and Recall. The two metrics are also calculated considering a one-versus-all approach. The F-score for the multi-class case can be calculated using its average per class values [39]:

$$F\text{-score} = 2 \times \frac{\overline{Precision} \times \overline{Recall}}{\overline{Precision} + \overline{Recall}}$$

- The OA consists of the number of TP divided by the total amount of observations. Considering c as the label for the different classes present in a target class, OA is given by the following formula:

$$OA = \frac{\sum_c TP_c}{\sum_c (TP_c + FP_c)}$$

The comparison of the performance of AL frameworks is based on its data selection and augmentation efficacy. Specifically, an efficient data selection/generation policy allows the production of classifiers with high performance on unseen data while using as least non-artificial training data as possible. To measure the performance of the different AL setups, we follow the recommendations found in [42]. The performance of an AL setup will be compared using two AL-specific performance metrics:

- Area Under the Learning Curve (AULC). It is the sum of the classification performance over a validation/test set of the classifiers trained of all AL iterations. To facilitate the interpretability of this metric, the resulting AULC scores are fixed within the range $[0, 1]$ by dividing the AULC scores by the total amount of iterations (*i.e.*, the maximum performance area).
- Data Utilization Rate (DUR) [43]. Measures the percentage of training data required to reach a given performance threshold, as a ratio of the percentage of training data required by the baseline framework. This metric is also presented as a percentage of the total amount of training data, without making it relative to the baseline framework. The DUR metric is measured at 45 different performance thresholds, ranging between $[0.10, 1.00]$ at a 0.02 step.

383 4.4. Experimental Procedure

384

385 The evaluation of different active learners in a live setting is generally ex-
 386 pensive, time-consuming and prone to human error. Instead, a common prac-
 387 tice is to compare them in an offline environment using labeled datasets [44].
 388 In this scenario, since the dataset is already labeled, the annotation process
 389 is done at zero cost. Figure 5 depicts the experiment designed for one dataset
 390 over a single run.

391 A single run starts with the splitting of a preprocessed dataset in 5 dif-
 392 ferent partitions, stratified according to the class frequencies of the target
 393 variable using the K-fold Cross Validation method. During this run, an ac-
 394 tive learner or classifier is trained 5 times using a different partition as the
 395 Test set each time. For each training process, a Validation set containing 25%
 396 of the subset is created and is used to measure the data selection efficiency
 397 (*i.e.*, AULC and DUR using the classification performance metrics, specific
 398 to AL). Therefore, for a single training procedure, 20% of the original dataset
 399 is used as the Validation set, 20% is used as the Test set and 60% is used as
 400 the Train set. The AL simulations and the classifiers' training occur within
 401 the Train set. However, the classifiers used to find the maximum performance
 402 classification scores are trained over the full Train set. The AL simulations
 403 are run over a maximum of 50 iterations (including the initialization step),
 404 adding 1.6% of the training set each time (*i.e.*, all AL simulations use less
 405 than 80% of the Train set). Once the training phase is completed, the Test
 406 set classification scores are calculated using the trained classifiers. For the
 407 case of AL, the classifier with the optimal Validation set score is used to
 408 estimate the AL's optimal classification performance over unseen data.

409 The process shown in Figure 5 is repeated over 3 runs using different
 410 random seeds over the 10 different datasets collected. The final scores of
 411 each AL configuration and classifier correspond to the average of the 3 runs
 412 and 5-fold Cross Validation estimations (*i.e.*, the mean score of 15 fits, across
 413 10 datasets).

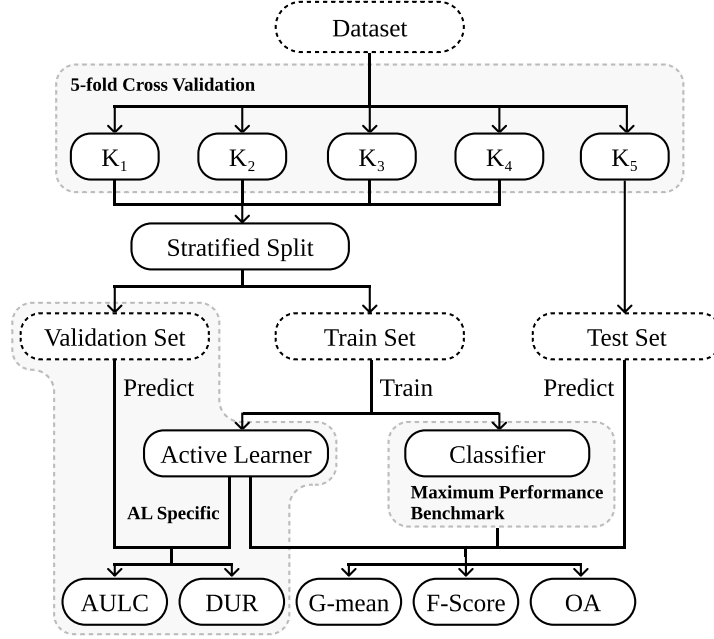


Figure 5: Experimental procedure flowchart. The preprocessed datasets are split into five folds. One of the folds is used to test the best found classifiers using AL and the classifiers trained using the entire training dataset (containing the remaining folds). The training set is used to run both the AL simulations as well as train the normal classifiers. The validation set is used to measure AL-specific performance metrics over each iteration. We use different subsets for overall classification performance and AL-specific performance to avoid data leakage.

414 The hyperparameters defined for the AL frameworks, Classifiers and Gen-
415 erators are shown in Table 2. In the Generators table, we distinguish the
416 G-SMOTE algorithm working as a normal oversampling method from G-
417 SMOTE-AUGM, which performs generates additional artificial data on top
418 of the usual oversampling mechanism. Since the G-SMOTE-AUGM method
419 is intended to be used with varying parameter values (via within-iteration
420 parameter tuning), the parameters were defined as a list of various possible
421 values.

Active Learners	Hyperparameters	Inputs
Standard	# initial obs.	1.6%
	# additional obs. per iteration	1.6%
	max. iterations + initialization	50
	evaluation metrics	G-mean, F-score, OA
	selection strategy	Random, Entropy, Breaking Ties
	within-iteration param. tuning	None
	generator	None
	classifier	DT, LR, KNN, RF
Oversampling	generator	G-SMOTE
Proposed	generator	G-SMOTE-AUGM
	within-iteration param. tuning	Grid Search K-fold CV
Classifier		
DT	min. samples split	2
	criterion	gini
LR	maximum iterations	100
	multi class	One-vs-All
	solver	liblinear
KNN	penalty	L2 (Ridge)
	# neighbors	5
	weights	uniform
	metric	euclidean
RF	min. samples split	2
	# estimators	100
	criterion	gini
Generator		
G-SMOTE	# neighbors	4
	deformation factor	0.5
	truncation factor	0.5
G-SMOTE-AUGM	# neighbors	3, 4, 5
	deformation factor	0.5
	truncation factor	0.5
	augmentation factor	[1.1, 2.0] at 0.1 step

Table 2: Hyperparameter definition for the active learners, classifiers and generators used in the experiment.

422 4.5. Software Implementation

423

424 The experiment was implemented using the Python programming lan-
425 guage, along with the Python libraries Scikit-Learn [45], Imbalanced-Learn [46],
426 Geometric-SMOTE [15], Research-Learn and ML-Research libraries. All
427 functions, algorithms, experiments and results are provided in the GitHub
428 repository of the project.

429 5. Results & Discussion

430

431 In a multiple dataset experiment, the analysis of results should not rely
432 uniquely on the average performance scores across datasets. The domain of
433 application and fluctuations of performance scores between datasets make
434 the analysis of these averaged results less accurate. Instead, it is generally
435 recommended the use of the mean ranking scores to extend the analysis [47].
436 Since mean performance scores are still intuitive to interpret, we will present
437 and discuss both results. The rank values are assigned based on the mean
438 scores of 3 different runs of 5-fold Cross Validation (15 performance estima-
439 tions per dataset) for each combination of dataset, AL configuration, classifier
440 and performance metric.

441 5.1. Results

442

443 The average ranking of the AULC estimations of AL methods are shown
444 in Table 3. The proposed method almost always improves AL performance
445 and ensures higher data selection efficiency.

Classifier	Evaluation Metric	Standard	Oversampling	Proposed
DT	Accuracy	2.50 ± 0.81	2.20 ± 0.40	1.30 ± 0.64
DT	F-score	2.50 ± 0.81	2.10 ± 0.30	1.40 ± 0.80
DT	G-mean	2.70 ± 0.64	2.00 ± 0.45	1.30 ± 0.64
KNN	Accuracy	2.40 ± 0.80	1.90 ± 0.54	1.70 ± 0.90
KNN	F-score	2.60 ± 0.66	1.80 ± 0.40	1.60 ± 0.92
KNN	G-mean	2.80 ± 0.40	1.70 ± 0.46	1.50 ± 0.81
LR	Accuracy	2.60 ± 0.66	2.10 ± 0.54	1.30 ± 0.64
LR	F-score	2.80 ± 0.40	2.00 ± 0.45	1.20 ± 0.60
LR	G-mean	2.80 ± 0.40	2.00 ± 0.45	1.20 ± 0.60
RF	Accuracy	2.60 ± 0.66	1.90 ± 0.54	1.50 ± 0.81
RF	F-score	2.60 ± 0.66	2.00 ± 0.45	1.40 ± 0.80
RF	G-mean	2.80 ± 0.40	1.60 ± 0.49	1.60 ± 0.80

Table 3: Mean rankings of the AULC metric over the different datasets (10), folds (5) and runs (3) used in the experiment. The proposed method always improves the results of the original framework and on average almost always improves the results of the oversampling framework.

446 Table 4 shows the average AULC scores, grouped by classifier, Evalua-
447 tion Metric and AL framework. The variation in performance across active
448 learners is consistent with the mean rankings found in Table 3, while showing
449 significant AULC score differences between the proposed AL method and the
450 oversampling AL method.

451 The average DUR scores were calculated for various G-mean thresholds,
452 varying between 0.1 and 1.0 at a 0.02 step (45 different thresholds in total).
453 Table 5 shows the results obtained for these scores starting from a G-mean
454 score of 0.6 and was filtered to show only the thresholds ending with 0 or 6.
455 In most cases, the proposed method reduces the amount of data annotation
456 required to reach each G-mean score threshold.

457 The DUR scores relative to the Standard AL method are shown in Fig-
458 ure 6. A DUR below 1 means that the Proposed/Oversampling method
459 requires less data than the Standard AL method to reach the same perfor-
460 mance threshold. For example, running an AL simulation using the KNN
461 classifier requires 69.6% of the amount of data required by the Standard AL
462 method using the same classifier to reach an F-Score of 0.62 (*i.e.*, requires
463 30.4% less data).

464 The mean optimal classification scores of AL methods and Classifiers

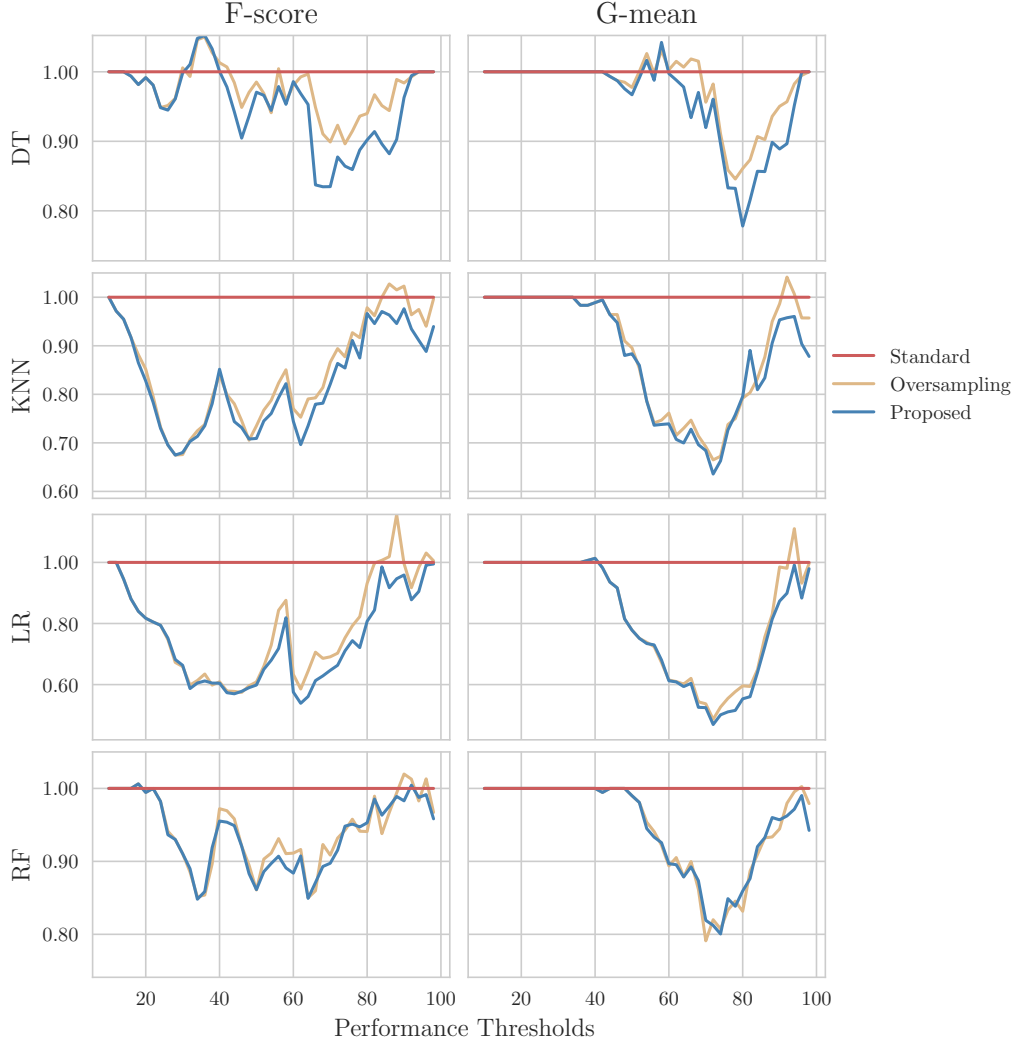


Figure 6: Mean data utilization rates. The y-axis shows the percentage of data (relative to the baseline AL framework) required to reach the different performance thresholds.

Classifier	Evaluation Metric	Standard	Oversampling	Proposed
DT	Accuracy	0.733 ± 0.092	0.732 ± 0.087	0.740 ± 0.087
DT	F-score	0.695 ± 0.088	0.698 ± 0.090	0.705 ± 0.092
DT	G-mean	0.804 ± 0.065	0.811 ± 0.060	0.816 ± 0.062
KNN	Accuracy	0.816 ± 0.091	0.818 ± 0.088	0.822 ± 0.091
KNN	F-score	0.775 ± 0.102	0.784 ± 0.108	0.788 ± 0.111
KNN	G-mean	0.852 ± 0.084	0.866 ± 0.072	0.869 ± 0.074
LR	Accuracy	0.802 ± 0.091	0.812 ± 0.088	0.821 ± 0.086
LR	F-score	0.749 ± 0.112	0.773 ± 0.116	0.784 ± 0.115
LR	G-mean	0.839 ± 0.093	0.870 ± 0.065	0.875 ± 0.064
RF	Accuracy	0.861 ± 0.076	0.861 ± 0.075	0.862 ± 0.077
RF	F-score	0.823 ± 0.105	0.827 ± 0.105	0.829 ± 0.105
RF	G-mean	0.886 ± 0.077	0.895 ± 0.063	0.895 ± 0.065

Table 4: Average AULC of each AL configuration tested. Each AULC score is calculated using the performance scores of each iteration in the validation set. By the end of the iterative process, each AL configuration used a maximum of 80% instances of the 60% instances that compose the training sets (*i.e.*, 48% of the entire preprocessed dataset).

(fully labeled training set, without AL) is shown in Table 6. The proposed AL method produces classifiers that are almost always able to outperform classifiers using the full training set (*i.e.*, the ones labeled as MP).

5.2. Statistical Analysis

When checking for statistical significance in a multiple dataset context it is important to account for the multiple comparison problem. Consequently, our statistical analysis focuses on the recommendations found in [47]. Overall, we perform 3 statistical tests. The Friedman test [48] is used to understand whether there is a statistically significant difference in performance between the 3 AL frameworks. As post hoc analysis, the Wilcoxon signed-rank test [49] was used to check for statistical significance between the performance of the proposed AL method and the oversampling AL method across datasets. As a second post hoc analysis, the Holm-Bonferroni [50] method was used to check for statistical significance between the methods using data generators and the Standard AL framework across classifiers and evaluation metrics.

G-mean Score	Classifier	Standard	Oversampling	Proposed
0.60	DT	3.2%	3.1%	3.2%
0.60	KNN	3.6%	2.6%	2.5%
0.60	LR	3.9%	2.2%	2.2%
0.60	RF	2.4%	2.1%	2.1%
0.66	DT	4.6%	4.6%	4.2%
0.66	KNN	4.9%	3.7%	3.5%
0.66	LR	5.7%	3.2%	3.1%
0.66	RF	3.0%	2.8%	2.7%
0.70	DT	6.6%	6.1%	5.8%
0.70	KNN	8.5%	5.0%	4.7%
0.70	LR	9.5%	4.6%	4.3%
0.70	RF	4.5%	3.2%	3.3%
0.76	DT	16.5%	13.0%	12.7%
0.76	KNN	17.8%	9.7%	9.0%
0.76	LR	16.6%	10.0%	7.8%
0.76	RF	10.1%	5.5%	5.5%
0.80	DT	36.1%	30.4%	27.1%
0.80	KNN	22.7%	18.0%	17.8%
0.80	LR	25.2%	16.0%	14.2%
0.80	RF	15.5%	9.0%	9.5%
0.86	DT	60.5%	56.7%	54.5%
0.86	KNN	39.9%	37.0%	37.8%
0.86	LR	32.6%	27.5%	27.0%
0.86	RF	28.0%	25.7%	25.7%
0.90	DT	72.5%	70.7%	67.8%
0.90	KNN	49.9%	50.3%	49.3%
0.90	LR	52.5%	53.8%	49.3%
0.90	RF	44.6%	42.6%	43.5%
0.96	DT	100.0%	99.5%	100.0%
0.96	KNN	79.4%	75.6%	71.6%
0.96	LR	87.5%	83.1%	79.8%
0.96	RF	63.6%	64.2%	63.1%

Table 5: Mean data utilization of AL algorithms, as a percentage of the training set.

Classifier	Evaluation Metric	MP	Standard	Oversampling	Proposed
DT	Accuracy	0.809 ± 0.086	0.802 ± 0.089	0.806 ± 0.089	0.812 ± 0.087
DT	F-score	0.774 ± 0.107	0.772 ± 0.096	0.775 ± 0.101	0.781 ± 0.103
DT	G-mean	0.853 ± 0.081	0.854 ± 0.069	0.860 ± 0.067	0.864 ± 0.068
KNN	Accuracy	0.882 ± 0.085	0.883 ± 0.087	0.877 ± 0.087	0.881 ± 0.093
KNN	F-score	0.848 ± 0.116	0.849 ± 0.115	0.847 ± 0.118	0.852 ± 0.121
KNN	G-mean	0.896 ± 0.094	0.899 ± 0.090	0.904 ± 0.078	0.907 ± 0.080
LR	Accuracy	0.855 ± 0.074	0.870 ± 0.073	0.858 ± 0.077	0.870 ± 0.076
LR	F-score	0.812 ± 0.113	0.835 ± 0.105	0.825 ± 0.106	0.838 ± 0.106
LR	G-mean	0.875 ± 0.099	0.895 ± 0.075	0.899 ± 0.059	0.907 ± 0.059
RF	Accuracy	0.897 ± 0.080	0.905 ± 0.078	0.904 ± 0.078	0.906 ± 0.077
RF	F-score	0.867 ± 0.107	0.877 ± 0.103	0.875 ± 0.108	0.877 ± 0.108
RF	G-mean	0.911 ± 0.081	0.917 ± 0.078	0.923 ± 0.067	0.925 ± 0.065

Table 6: Optimal classification scores. The Maximum Performance (MP) classification scores are calculated using classifiers trained using the entire training set.

482 Table 7 contains the *p-values* obtained with the Friedman test. The
483 difference in performance across AL frameworks is statistically significant at
484 a level of $\alpha = 0.05$ regardless of the classifier or evaluation metric being
485 considered.

Classifier	Evaluation Metric	p-value	Significance
DT	Accuracy	2.1e-17	True
DT	F-score	2.5e-24	True
DT	G-mean	2.8e-16	True
KNN	Accuracy	1.1e-46	True
KNN	F-score	1.8e-66	True
KNN	G-mean	6.4e-42	True
LR	Accuracy	9.9e-59	True
LR	F-score	2.0e-76	True
LR	G-mean	2.2e-59	True
RF	Accuracy	5.7e-42	True
RF	F-score	4.6e-55	True
RF	G-mean	1.3e-38	True

Table 7: Results for Friedman test. Statistical significance is tested at a level of $\alpha = 0.05$. The null hypothesis is that there is no difference in the classification outcome across oversamplers.

Table 8 contains the *p-values* obtained with the Wilcoxon signed-rank test. The proposed method was able to outperform both the standard AL framework, as well as the AL framework using a normal oversampling policy proposed in [18] with statistical significance in 9 out of 10 datasets.

The *p-values* shown in Table 9 refer to the results of the Holm-Bonferroni test. The proposed method’s superior performance was statistically significant for any combination of classifier and evaluation metric. Simultaneously, the proposed method established statistical significance in the 3 scenarios where the oversampling AL method failed to do so.

5.3. Discussion

In this paper we study the application of data augmentation methods through the modification of the standard AL framework. This is done to further reduce the amount of labeled data required to produce a reliable classifier, at the expense of artificial data generation.

In Table 3 we found that the proposed method was able to outperform the Standard AL framework in all scenarios. The mean rankings are consistent with the mean AULC scores found in Table 4, while showing significant performance differences between the proposed method and both the standard

Dataset	Oversampling	Standard
Baseball	5.0e-01	3.4e-01
Gas Drift	3.7e-26	4.6e-57
Image Segmentation	9.6e-18	2.1e-44
Japanese Vowels	2.4e-09	1.6e-32
Mfeat Zernike	1.2e-12	9.5e-40
Mice Protein	6.5e-32	1.5e-61
Pendigits	5.0e-18	2.3e-45
Texture	1.5e-22	6.7e-57
Vehicle	7.4e-11	7.9e-13
Waveform	8.9e-08	2.6e-02

Table 8: Adjusted p-values using the Wilcoxon signed-rank method. Bold values are statistically significant at a level of $\alpha = 0.05$. The null hypothesis is that the performance of the proposed framework is similar to that of the oversampling or standard framework.

Classifier	Evaluation Metric	Oversampling	Proposed
DT	Accuracy	4.5e-05	1.6e-10
DT	F-score	1.9e-07	2.7e-10
DT	G-mean	2.5e-06	3.1e-09
KNN	Accuracy	5.5e-02	1.1e-05
KNN	F-score	6.7e-11	6.3e-14
KNN	G-mean	8.3e-06	1.3e-07
LR	Accuracy	8.1e-02	3.4e-06
LR	F-score	7.1e-06	2.0e-20
LR	G-mean	2.2e-07	1.1e-11
RF	Accuracy	2.0e-01	2.8e-02
RF	F-score	2.2e-05	8.1e-07
RF	G-mean	2.0e-04	2.0e-04

Table 9: Adjusted p-values using the Holm-Bonferroni method. Bold values are statistically significant at a level of $\alpha = 0.05$. The null hypothesis is that the Oversampling or Proposed method does not perform better than the control method (Standard AL framework).

505 and oversampling methods. The Friedman test in Table 7 showed that the
506 difference in the performance of these AL frameworks is statistically signifi-
507 cant, regardless of the classifier or performance metric being used.

508 The proposed method showed more consistent data utilization require-
509 ments to most of the assessed G-mean score thresholds when compared to
510 the remaining AL methods, as seen in Table 5. For example, to reach a
511 G-mean Score of 0.9 using the KNN and LR classifiers, the average amount
512 of data required with the Oversampling AL approach increased when com-
513 pared to the Standard approach. However, the proposed method was able to
514 decrease the amount of data required in both situations. The robustness of
515 the Proposed method is clearer in Figure 6. In most cases, this method was
516 able outperform the Oversampling method. At the same time, the proposed
517 method also addresses inconsistencies in situations where the Oversampling
518 method was unable to outperform the standard method.

519 The statistical analyses found in Tables 8 and 9 showed that the pro-
520 posed method’s superiority was statistically significant in all datasets except
521 one (Baseball) and established statistical significance when compared to the
522 Standard AL method for all combinations of classifier and performance met-
523 ric, including when the Oversampling AL method failed to do so. These
524 results show that the Proposed method increased the reliability of the new
525 AL framework and improved the quality of the final classifier while using less
526 data.

527 Even though it was not the core purpose of this study, we found that
528 the method proposed AL approach consistently outperformed the maximum
529 performance threshold. Specifically, in Table 6, the performance of the classi-
530 fiers originating from the proposed method was able to outperform classifiers
531 trained using the full training dataset in all 12 scenarios except one. This
532 suggests that the selection of a meaningful training subset training dataset
533 paired with data augmentation not only matches the classification perfor-
534 mance of ML algorithms, as it also improves them. Even in a setting with
535 fully labeled training data, the proposed method may be used as preprocess-
536 ing method to further optimize classification performance.

537 This study discussed the effect of data augmentation within the AL frame-
538 work, along with the exploration of optimal augmentation methods within
539 AL iterations. However, the conceptual nature of this study implies some
540 limitations. Specifically, the large amount of experiments required to test
541 the method’s efficacy, along with the limited computational power available,
542 led to a limited exploration of the grid search’s potential. Future work should

focus into understanding how the usage of a more comprehensive parameter tuning approach improves the quality of the AL method. In addition, the proposed method was not able to outperform the standard AL method in 100% of scenarios. The exploration of other, more complex, data augmentation techniques might further improve its performance through the production of more meaningful training observations. Specifically, in this study we assume that all datasets used follow a manifold, allowing the usage of G-SMOTE as a data augmentation approach. However, this method cannot be used into more complex, non-euclidean spaces. In this scenario, the usage of G-SMOTE is not valid and might lead to the production of noisy data. Deep Learning-based data augmentation techniques are able to address this limitation and improve the overall quality of the artificial data being generated. We also found significant standard errors throughout our experimental results (see Subsection 5.1), which is consistent with the findings in [18, 42]. This suggests that the usage of more robust generators did not decrease the standard error of AL performance. Instead, AL’s performance variability is likely dependent on the quality of its initialization.

6. Conclusion

The ability of training ML classifiers is usually limited to the availability of labeled data. However, manually labeling data is often expensive, which makes the usage of AL particularly appealing to select the most informative observations and reduce the amount of required labeled data. On the other hand, the introduction of data variability in the training dataset can also be done via data augmentation. However, most, if not all, AL configurations using some form data augmentation are domain and/or task specific. These methods typically explore deep learning approaches on both classification and data augmentation. Consequently, they may not be applicable for other classification tasks or when the available computational power is insufficient.

In this paper, we proposed a domain-agnostic AL framework that implements Data Augmentation and hyperparameter tuning. We found that a heuristic Data Augmentation algorithm is sufficient to improve the data selection efficiency in AL. Specifically, the data augmentation method used almost always increased AL performance, regardless of the target goal (*i.e.*, optimizing classification or data selection efficiency). The usage of data augmentation reduced the **number** of iterations required to train a classifier with

579 a performance as good as (or better than) classifiers trained with the entire
580 training dataset (*i.e.*, without using AL). In addition, the proposed method
581 reduced the size of the training dataset, which is expanded with artificial
582 data.

583 With this AL configuration, data selection in AL iterations aim towards
584 observations that optimize the quality of the artificial data produced. The
585 substitution of less informative labeled data with artificial data is especially
586 useful in this context, since it allows the reduction of some of the user interac-
587 tion necessary to reach a sufficiently informative dataset. In order to further
588 improve the proposed method future work will (1) focus on the development
589 of methods with varying data augmentation policies depending on the differ-
590 ent input space regions, (2) develop augmentation-sensitive query functions
591 capable of avoiding the unnecessary selection of similar observations from the
592 unlabeled dataset and (3) better understand the gap between heuristic/input
593 space data augmentation techniques and neural network/feature space data
594 augmentation techniques in an AL context.

595 **Declarations**

596 *Funding*

597 This research was supported by three research grants of the Portuguese
598 Foundation for Science and Technology (“Fundação para a Ciência e a Tec-
599 nologia”), references SFRH/BD/151473/2021, DSAIPA/DS/0116/2019 and
600 PCIF/SSI/0102/2017.

601 *Code availability*

602 The analyses and source code is available at [github.com/joaopfonseca/ml-](https://github.com/joaopfonseca/ml-research)
603 research.

604 **References**

- 605 [1] V. Nath, D. Yang, B. A. Landman, D. Xu, H. R. Roth, Diminishing
606 uncertainty within the training pool: Active learning for medical image
607 segmentation, *IEEE Transactions on Medical Imaging* 40 (10) (2021)
608 2534–2547.
- 609 [2] Y. Sverchkov, M. Craven, A review of active learning approaches to
610 experimental design for uncovering biological networks, *PLoS Compu-*
611 *tational Biology* 13 (2017) e1005466.

- 612 [3] X. Li, D. Kuang, C. X. Ling, Active learning for hierarchical text classification,
613 Lecture Notes in Computer Science (including subseries Lecture
614 Notes in Artificial Intelligence and Lecture Notes in Bioinformatics) 7301
615 LNAI (2012) 14–25.
- 616 [4] Y. Li, J. Yin, L. Chen, Seal: Semisupervised adversarial active learning
617 on attributed graphs, IEEE Transactions on Neural Networks and
618 Learning Systems 32 (7) (2021) 3136–3147.
- 619 [5] O. Siméoni, M. Budnik, Y. Avrithis, G. Gravier, Rethinking deep active
620 learning: Using unlabeled data at model training, Proceedings -
621 International Conference on Pattern Recognition (2020) 1220–1227.
- 622 [6] J. E. Van Engelen, H. H. Hoos, A survey on semi-supervised learning,
623 Machine Learning 109 (2) (2020) 373–440.
- 624 [7] O. Sener, S. Savarese, Active learning for convolutional neural networks:
625 A core-set approach, in: International Conference on Learning Representations,
626 2018.
- 627 [8] Y. Leng, X. Xu, G. Qi, Combining active learning and semi-supervised
628 learning to construct svm classifier, Knowledge-Based Systems 44 (2013)
629 121–131.
- 630 [9] H. Yu, X. Yang, S. Zheng, C. Sun, Active learning from imbalanced
631 data: A solution of online weighted extreme learning machine, IEEE
632 Transactions on Neural Networks and Learning Systems 30 (2019) 1088–
633 1103.
- 634 [10] W. Zong, G.-B. Huang, Y. Chen, Weighted extreme learning machine
635 for imbalance learning, Neurocomputing 101 (2013) 229–242.
- 636 [11] J. Qin, C. Wang, Q. Zou, Y. Sun, B. Chen, Active learning with extreme
637 learning machine for online imbalanced multiclass classification,
638 Knowledge-Based Systems 231 (2021) 107385.
- 639 [12] W. Liu, H. Zhang, Z. Ding, Q. Liu, C. Zhu, A comprehensive active
640 learning method for multiclass imbalanced data streams with concept
641 drift, Knowledge-Based Systems 215 (2021) 106778.

- 642 [13] A. Tharwat, W. Schenck, Balancing exploration and exploitation: A
643 novel active learner for imbalanced data, *Knowledge-Based Systems* 210
644 (2020) 106500.
- 645 [14] Y.-Y. Kim, K. Song, J. Jang, I.-c. Moon, Lada: Look-ahead data ac-
646 quisition via augmentation for deep active learning, *Advances in Neural*
647 *Information Processing Systems* 34 (2021).
- 648 [15] G. Douzas, F. Bacao, Geometric SMOTE a geometrically enhanced
649 drop-in replacement for SMOTE, *Information Sciences* 501 (2019) 118–
650 135.
- 651 [16] J. Katz-Samuels, J. Zhang, L. Jain, K. Jamieson, Improved algorithms
652 for agnostic pool-based active classification, in: *International Conference*
653 *on Machine Learning*, PMLR, 2021, pp. 5334–5344.
- 654 [17] T. Su, S. Zhang, T. Liu, Multi-spectral image classification based on an
655 object-based active learning approach, *Remote Sensing* 12 (2020) 504.
- 656 [18] J. Fonseca, G. Douzas, F. Bacao, Increasing the Effectiveness of Active
657 Learning: Introducing Artificial Data Generation in Active Learning
658 for Land Use/Land Cover Classification, *Remote Sensing* 2021, Vol. 13,
659 Page 2619 13 (13) (2021) 2619.
- 660 [19] D. Yoo, I. S. Kweon, Learning loss for active learning, in: *Proceedings*
661 *of the IEEE/CVF Conference on Computer Vision and Pattern Recog-*
662 *nition*, 2019, pp. 93–102.
- 663 [20] H. H. Aghdam, A. Gonzalez-Garcia, A. Lopez, J. Weijer, Active learn-
664 ing for deep detection neural networks, in: *Proceedings of the IEEE*
665 *International Conference on Computer Vision*, Vol. 2019-Octob, 2019,
666 pp. 3671–3679.
- 667 [21] B. Gu, Z. Zhai, C. Deng, H. Huang, Efficient active learning by querying
668 discriminative and representative samples and fully exploiting unlabeled
669 data, *IEEE Transactions on Neural Networks and Learning Systems* 32
670 (2021) 4111–4122.
- 671 [22] P. Kumar, A. Gupta, Active learning query strategies for classification,
672 regression, and clustering: A survey, *Journal of Computer Science and*
673 *Technology* 2020 35:4 35 (2020) 913–945.

- [23] Y. Fu, X. Zhu, B. Li, A survey on instance selection for active learning, Knowledge and information systems 35 (2) (2013) 249–283.
- [24] S. Behpour, K. M. Kitani, B. D. Ziebart, Ada: Adversarial data augmentation for object detection, Proceedings - 2019 IEEE Winter Conference on Applications of Computer Vision, WACV 2019 (2019) 1243–1252.
- [25] Z. Zhong, L. Zheng, G. Kang, S. Li, Y. Yang, Random erasing data augmentation, in: Proceedings of the AAAI Conference on Artificial Intelligence, Vol. 34, 2020, pp. 13001–13008.
- [26] T. DeVries, G. W. Taylor, Dataset augmentation in feature space, in: 5th International Conference on Learning Representations, ICLR 2017 - Workshop Track Proceedings, International Conference on Learning Representations, ICLR, 2017.
- [27] C. Shorten, T. M. Khoshgoftaar, A survey on image data augmentation for deep learning, Journal of Big Data 6 (1) (2019) 1–48.
- [28] O. Kashefi, R. Hwa, Quantifying the evaluation of heuristic methods for textual data augmentation, in: Proceedings of the Sixth Workshop on Noisy User-generated Text (W-NUT 2020), Association for Computational Linguistics, Online, 2020, pp. 200–208.
- [29] T. Tran, T.-T. Do, I. Reid, G. Carneiro, Bayesian generative active deep learning, in: International Conference on Machine Learning, PMLR, 2019, pp. 6295–6304.
- [30] S. Sinha, S. Ebrahimi, T. Darrell, Variational adversarial active learning, in: Proceedings of the IEEE/CVF International Conference on Computer Vision, 2019, pp. 5972–5981.
- [31] K. Kim, D. Park, K. I. Kim, S. Y. Chun, Task-aware variational adversarial active learning, in: Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2021, pp. 8166–8175.
- [32] Y. Ma, S. Lu, E. Xu, T. Yu, L. Zhou, Combining active learning and data augmentation for image classification, in: Proceedings of the 2020 3rd International Conference on Big Data Technologies, 2020, pp. 58–62.

- [33] H. Quteineh, S. Samothrakis, R. Sutcliffe, Textual data augmentation for efficient active learning on tiny datasets, in: Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP), 2020, pp. 7400–7410.
- [34] Q. Li, Z. Huang, Y. Dou, Z. Zhang, A framework of data augmentation while active learning for chinese named entity recognition, in: International Conference on Knowledge Science, Engineering and Management, Springer, 2021, pp. 88–100.
- [35] C. Wu, The decision tree approach to classification., Purdue University, 1975.
- [36] T. Cover, P. Hart, Nearest neighbor pattern classification, IEEE Transactions on Information Theory 13 (1) (1967) 21–27.
- [37] T. K. Ho, Random decision forests, in: Proceedings of the Third International Conference on Document Analysis and Recognition (Volume 1) - Volume 1, ICDAR '95, IEEE Computer Society, USA, 1995, p. 278.
- [38] J. A. Nelder, R. W. Wedderburn, Generalized linear models, Journal of the Royal Statistical Society: Series A (General) 135 (3) (1972) 370–384.
- [39] L. A. Jeni, J. F. Cohn, F. De La Torre, Facing imbalanced data - Recommendations for the use of performance metrics, in: Proceedings - 2013 Humaine Association Conference on Affective Computing and Intelligent Interaction, ACII 2013, 2013, pp. 245–251.
- [40] M. Fatourechi, R. K. Ward, S. G. Mason, J. Huggins, A. Schloegl, G. E. Birch, Comparison of evaluation metrics in classification applications with imbalanced datasets, in: 2008 seventh international conference on machine learning and applications, IEEE, 2008, pp. 777–782.
- [41] M. Kubat, S. Matwin, et al., Addressing the curse of imbalanced training sets: one-sided selection, in: Icml, Vol. 97, Citeseer, 1997, pp. 179–186.
- [42] D. Kottke, A. Calma, D. Huseljic, G. Kremlpl, B. Sick, Challenges of reliable, realistic and comparable active learning evaluation, in: CEUR Workshop Proceedings, Vol. 1924, 2017, pp. 2–14.

- 734 [43] T. Reitmaier, B. Sick, Let us know your decision: Pool-based active
735 training of a generative classifier with the selection strategy 4ds, Infor-
736 mation Sciences 230 (2013) 106–131.
- 737 [44] J.-F. Kagy, T. Kayadelen, J. Ma, A. Rostamizadeh, J. Strnadova, The
738 practical challenges of active learning: Lessons learned from live exper-
739 imentation, arXiv preprint arXiv:1907.00038 (6 2019).
- 740 [45] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion,
741 O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vander-
742 plas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, É. Duchesnay,
743 Scikit-learn: Machine Learning in Python, Journal of Machine Learning
744 Research 12 (Oct) (2011) 2825–2830.
- 745 [46] G. Lemaître, F. Nogueira, C. K. Aridas, Imbalanced-learn: A python
746 toolbox to tackle the curse of imbalanced datasets in machine learning,
747 Journal of Machine Learning Research 18 (17) (2017) 1–5.
- 748 [47] J. Demšar, Statistical comparisons of classifiers over multiple data sets,
749 Journal of Machine Learning Research 7 (2006) 1–30.
- 750 [48] M. Friedman, The use of ranks to avoid the assumption of normality
751 implicit in the analysis of variance, Journal of the american statistical
752 association 32 (200) (1937) 675–701.
- 753 [49] F. Wilcoxon, Individual Comparisons by Ranking Methods, Biometrics
754 Bulletin 1 (6) (1945) 80.
- 755 [50] S. Holm, A simple sequentially rejective multiple test procedure, Scan-
756 dinavian journal of statistics (1979) 65–70.