

Tabular synthetic data generation: A literature review

Joao Fonseca^{1*}, Fernando Bacao¹

¹NOVA Information Management School, Universidade Nova de Lisboa

*Corresponding Author

Postal Address: NOVA Information Management School, Campus de Campolide, 1070-312 Lisboa, Portugal

Telephone: +351 21 382 8610

The generation of synthetic data can be used for anonymization, regularization, oversampling, semi-supervised learning, self-supervised learning and various other tasks. The wide range of applications of these mechanisms motivated the development of new algorithms specialized in generating data for specific types of data and Machine Learning (ML) tasks. As a result, the analysis of the different types of generative models

1 Introduction

Synthetic data is obtained from a generative process based on properties of real data [1]. The generation of synthetic data is essential for various domains and tasks. For example, synthetic data is used as a form of regularizing neural networks (*i.e.*, data augmentation) [CITATION]. One form of anonymizing datasets is via the production of synthetic observations (*i.e.*, synthetic data generation) [CITATION]. In settings where only a small portion of training data is labeled, some techniques generate artificial data using both labeled and unlabeled data with a modified loss function to train neural networks (*i.e.*, semi-supervised learning) [2]. In imbalanced learning contexts, synthetic data can be used to balance the target classes' frequencies and reinforce the learning of minority classes (*i.e.*, oversampling) [3]. Some active learning frameworks use data generation to improve the quality of data selection and classifier training [4]. Other techniques employ data generation to produce deep neural networks without labeled data (*i.e.*, self-supervised learning) [5].

The breadth of these techniques span multiple domains, such as facial recognition [6], Land Use/Land Cover mapping [CITATION], medical image processing [CITATION], Natural Language Processing (NLP) [7] or credit card default prediction [8]. According to the domain and data type, the data generation techniques used may vary significantly. Generally speaking, some data generation mechanisms are specific to some domains, data types or tasks. For example, ... Most, if not all, of these techniques are applied on the input or output space.

However, there are various data generation techniques that are invariant to the task or data types used. These techniques can be either applied in the feature space [9] or in tabular datasets¹. On one hand,

¹Tabular data is a database structured in tabular form, composed of columns (features) and rows (observations) [10]

data generation in the feature space uses a generative model to learn a manifold, lower-dimensional abstraction over the input space [11], defined here as the feature space. At this level, any tabular data generation mechanism can be applied and reconstructed into the input space if necessary. On the other hand, synthetic data generation on tabular data can be applied to most problems. Although, the choice of generation mechanism is still dependant on (1) the importance of the relationships found between the different features, (2) the ML task developed and (3) the motivation for the generation of synthetic data. For example, when generating data to address an imbalanced learning problem (*i.e.*, oversampling), the relationships between the different features are not necessarily kept since the goal is to reinforce the learning of the minority class by redefining an ML classifier’s decision boundaries. If the goal is to anonymize a dataset, perform some type of descriptive task, or ensure a consistent model interpretability, these relationships need to be kept.

Depending on the context, evaluating the quality of the generated data is a complex task. For example, for image and time series data, perceptually small changes in the original data can lead to large changes in the euclidean distance [1, 12]. The evaluation of generative models typically account primarily for the performance in a specific task, since good performance in one criterion does not imply good performance on another [12]. However, in computationally intensive tasks it is often impracticable to search for the optimal configurations of generative models. To address this limitation, other evaluation methods have been proposed to assist in this evaluation, which can be distinguished into statistical divergence metrics and precision/recall metrics [13]. The relevant performance metrics found in the literature are discussed in Section 6.

1.1 Motivation, Scope and Contributions

This literature review focuses on generation mechanisms applied to tabular data and the different ML techniques where tabular synthetic data is used. In addition, we focus on the ML perspective of synthetic data, as opposed to the practical perspective. From a practical sense, synthetic data is used as a proxy of real data. It is assumed to be inaccessible, essential and a secondary asset for tasks like education, software development, or systems demonstrations [14].

We focus on data generation techniques in the tabular and feature space (*i.e.*, embedded inputs), given its breadth in scope. Related literature reviews are mostly focused on specific algorithmic or domain applications, with little to no emphasis on the core generative process. For this reason, these techniques often appear “sandboxed”, even though there is a significant overlap between them. There are some related reviews published since 2019. Assefa et al. [1] provides a general overview of synthetic data generation for time series data anonymization in the finance sector. Hernandez et al. [15] reviews data generation techniques for tabular health records anonymization. Raghunathan [16] reviews synthetic data anonymization techniques that preserve the statistical properties of a dataset. Nalepa et al. [17] reviews data augmentation techniques for brain-tumor segmentation. Bayer et al. [18] distinguishes augmentation techniques for text classification into feature and data space, while providing an extensive overview of augmentation methods within this domain. However, the taxonomy proposed and feature space augmentation methods are not necessarily specific to the domain. Shorten et al. [19], Chen et al. [20], Feng et al. [7] and Liu et al. [21] also review data augmentation techniques for text data. Yi et al. [22] review Generative Adversarial Network architectures for medical imaging. Wang et al. [23] reviews face data augmentation techniques. Shorten et al. [24] and Khosla et al. [25] discuss techniques for image data augmentation. Iwana et al. [26] and Wen et al. [27] also review time series data augmentation techniques. Zhao et al. [28] review data augmentation techniques for graph data. The analysis of related literature reviews ² is shown in Table 1.

²Results obtained using Google Scholar, limited to articles published since 2019, using the search

Table 1: Related literature reviews published since 2019.

Reference	Data type	ML problem	Domain	Observations
Assefa et al. [1]	—	Differential privacy	Finance	Analysis of applications, motivation and properties of synthetic data for anonymization.
Hernandez et al. [15]	Tabular	Differential privacy	Healthcare	Focus on GANs.
Raghunathan [16]	Tabular	Differential privacy	Statistics	Focus on general definitions such as differential privacy and statistical disclosure control.
Nalepa et al. [17]	Image	Segmentation	Medicine	Analysis of algorithmic applications on a 2018 brain-tumor segmentation challenge.
Bayer et al. [18]	Text	Classification	—	Distinguish 100 methods into 12 groups.
Shorten et al. [19]	Text	Deep Learning	—	General overview of text data augmentation.
Chen et al. [20]	Text	Few-shot Learning	—	Augmentation techniques for machine learning with limited data
Feng et al. [7]	Text	—	—	Overview of augmentation techniques and applications on NLP tasks.
Liu et al. [21]	Text	—	Various	Analysis of industry use cases of data augmentation in NLP. Emphasis on input level data augmentation.
Yi et al. [22]	Image	—	Medicine	Emphasis on GANs.
Wang et al. [23]	Image	Deep Learning	—	Regularization techniques using facial image data. Emphasis on Deep Learning generative models.
Shorten et al. [24]	Image	Deep Learning	—	Emphasis on data augmentation as a regularization technique.
Khosla et al. [25]	Image	—	—	Broad overview of image data augmentation. Emphasis on traditional approaches.
Iwana et al. [26]	Time series	Classification	—	Defined a taxonomy for time series data augmentation.
Wen et al. [27]	Time series	Various	—	Analysis of data augmentation methods for classification, anomaly detection and forecasting.
Zhao et al. [28]	Graph	Various	—	Graph data augmentation for supervised and self-supervised learning.
Khalifa et al. [29]	Image	—	Various	General overview of image data augmentation and relevant domains of application.

70 The different taxonomies established in the literature follow a similar philosophy, but vary in terminology
 71 and are often specific to the technique discussed. Regardless, it is possible to establish a broader taxonomy
 72 without giving up on specificity. This study provides a joint overview of the different data generation
 73 approaches, domains and ML techniques where data generation is being used, as well as a common
 74 taxonomy across domains. It extends the analyses found in these articles and uses the compiled knowledge
 75 to identify research gaps. We compare the strengths and weaknesses of the models developed within each
 76 of these fields. Finally, we identify possible future research directions to address some of the limitations
 77 found. The contributions of this paper are summarized below:

query ("synthetic data generation" OR "oversampling" OR "imbalanced learning" OR "data augmentation") AND ("literature review" OR "survey"). Retrieved on August 11th, 2022. More articles were added later whenever found relevant.

- Bridge different ML concepts using synthetic data generation in its core (Algorithmic applications + Review of the State-of-the-art).
- Propose a synthetic data generation/data augmentation taxonomy to resolve the ambiguity in the literature (Data augmentation taxonomy).
- Characterize all relevant data generation methods using the proposed taxonomy.
- Discuss the ML techniques in which synthetic data generation/data augmentation is used, beyond regularization and consolidate the current data generation mechanisms across the different techniques (Algorithmic Applications).
- Bring to light the key challenges of synthetic data generation and put forward possible research directions in the future.

1.2 Paper Organization

This paper is organized as follows: Section 2 defines and formalizes the different concepts, goals, trade-offs and motivations related to synthetic data generation. Section 3 establishes the taxonomy used to categorize all the methods described in the paper. Section ?? reviews synthetic data generation mechanisms in the feature space. Section ?? reviews synthetic data generation mechanisms in the input space. Section 5 describes the applications of synthetic data in ML methods. Section 6 reviews performance evaluation methods of synthetic data generation mechanisms. Section 7 summarizes the main findings and discusses limitations and possible research directions in the state-of-the-art. Section 8 presents the main conclusions drawn from this study.

2 Background

In this section we define basics concepts, common goals, trade-offs and motivations regarding the generation of synthetic data in ML. We define synthetic data generation as the production of observations using a generative model (regardless of its nature) that resemble naturally occurring observations within a certain domain. It requires access to either a training dataset, a generative process, or a data stream. However, additional requirements might be imposed depending on the ML task being developed. For example, to generate artificial data for regularization purposes in supervised learning (*i.e.*, data augmentation) the training dataset must be annotated [CITATION]. The generation of synthetic data for anonymization purposes assumes synthetic datasets to be different from the original data, while following the same statistical properties [CITATION]. Domain knowledge may also be necessary to encode specific relationships among features into the generative process.

2.1 Use Cases

The breach of sensitive information is an important barrier to the sharing of datasets, especially when it concerns personal information [30]. A common solution for this problem is the generation of synthetic data without identifiable information. Generally speaking, ML tasks that require data with sensitive information are not compromised when using synthetic data. The experiment conducted by Patki et al.

[31] using relational datasets showed that in 11 out 15 comparisons ($\approx 73\%$), practitioners performing predictive modelling tasks using fully synthetic datasets performed the same or better than those using the original dataset. This topic is discussed in Section 5.1.

A common problem in the training of deep neural networks are their capacity to generalize [32] (*i.e.*, reduce the difference in classification performance between known and unseen observations). Data augmentation is a common method to address this problem. The generation of synthetic observations increases the range of the possible input space used in the training phase, which reduces the performance difference between known and unseen observations. Although other regularization methods exist, data augmentation is a useful method since it does not affect the choice in the architecture of the ML classifier and does not exclude the usage of other regularization methods. In domains such as computer vision and NLP, data augmentation is also used to improve the robustness of models against adversarial attacks [33, 34]. These topics are discussed into higher detail in Section 5.2.

In supervised learning, synthetic data generation is often motivated by the need to balance target class distributions (*i.e.*, oversampling). Since most ML classifiers are designed to perform best with balanced datasets, defining an appropriate decision boundary to distinguish rare classes becomes difficult [35]. Although there are other approaches to address imbalanced learning, oversampling techniques are generally easier to implement since they do not involve modifications to the classifier. This topic is discussed into higher detail in Section 5.4.

In supervised learning projects where labeled data is not readily available, but can be labeled, an Active Learning (AL) method may be used to improve the labelling process. AL aims to reduce the cost of producing training datasets by finding the most informative observations to label and feed into the classifier [36]. In this case, the generation of synthetic data is particularly useful to reduce the amount of labelled data required for a successful ML project and its costs. A similar motivation applies to the case of few-shot learning: small datasets may be expanded with synthetic data [37]. These topics are discussed in Sections 5.5 and 5.6.

The two other techniques reliant on synthetic data generation is Semi-supervised and Self-supervised learning. The former leverages both labeled and unlabeled data in the training phase, simultaneously. Most of the methods in the literature apply perturbations on the training data as part of the training procedure [38]. Self-supervised learning is a technique used to train neural networks in the absence of labeled data. Both techniques use synthetic data generation as an internal procedure for most of these methods. These techniques are discussed in Sections 5.7 and 5.8.

2.2 Problem Formulation

The original dataset, $\mathcal{D} = \mathcal{D}_L \cup \mathcal{D}_U$, is a collection of real observations and is distinguished according to whether a target feature exists, $\mathcal{D}_L = ((x_i, y_i))_{i=1}^l$, or not, $\mathcal{D}_U = (x_i)_{i=1}^u$. All three datasets, \mathcal{D} , \mathcal{D}_L and \mathcal{D}_U consist of ordered collections with lengths $l + u$, l and u , respectively. Synthetic data generation is performed using a generator, $f_{gen}(x; \tau) = \tilde{x}$, where τ defines the generation policy (*i.e.*, its hyperparameters), $x \in \mathcal{D}$ is an observation and $\tilde{x} \in \mathcal{D}^s$ is a synthetic observation. Analogous to \mathcal{D} , the synthetic dataset, \mathcal{D}^s , is also distinguished according to whether there is an assignment of a target feature, $\mathcal{D}_L^s = ((\tilde{x}_j, \tilde{y}_j))_{j=1}^{l'}$, or not, $\mathcal{D}_U^s = (\tilde{x}_j)_{j=1}^{u'}$.

Depending on the ML task, it may be relevant to establish metrics to measure the quality of \mathcal{D}^s . In this case, a metric $f_{qual}(\mathcal{D}^s, \mathcal{D})$ is used to determine the level of similarity/dissimilarity between \mathcal{D} and \mathcal{D}^s . In addition, a performance metric to estimate the performance of a model on the objective task, f_{per} , may be

used to determine the appropriateness of a model with parameters θ , *i.e.*, f_θ . The generator’s goal is to generate \mathcal{D}^s with arbitrary length, given $\mathcal{D} \sim \mathbb{P}$ and $\mathcal{D}^s \sim \mathbb{P}^s$, such that $\mathbb{P}^s \approx \mathbb{P}$, $x_i \neq x_j \forall x_i \in \mathcal{D} \wedge x_j \in \mathcal{D}^s$. $f_{gen}(x; \tau)$ attempts to generate a \mathcal{D}^s that maximizes either f_{per} , f_{qual} , or a combination of both.

3 Data Generation Taxonomy

The taxonomy proposed in this paper is a compilation of different definitions found in the literature, along with other traits that vary among domains and generation techniques. Within image data studies, Shorten et al. [24] and Khalifa et al. [29] divide data augmentation techniques into “basic” or “classical” approaches and deep learning approaches. In both cases, the former refers to domain-specific generation techniques, while the latter may be applied to any type of data. Iwana et al. [26] proposes a time-series data augmentation taxonomy divided in four families: (1) Decomposition, (2) Pattern mixing, (3) Generative models and (4) Decomposition. With exception to generative models, the majority of the methods presented in the remaining families are well established and domain specific. Hernandez et al. [15] defines a taxonomy for synthetic tabular data generation approaches divided in three types of approaches: (1) Classical, (2) Deep learning and (3) Others. Most taxonomies found followed similar definitions with variations in terminology or distinction criteria. In addition, all taxonomies with categories defined as “basic”, “traditional” or “classical” use these to characterize domain-specific transformations.

Within the taxonomies found, none of them consider how a generation mechanism employs \mathcal{D} into the generation process or, if applicable, the training phase. However, it is important to understand whether a generation mechanism randomly selects x and a set of close neighbors, thus considering local information only, or considers the overall dataset or data distribution for the selection of x and/or generation of \tilde{x} . Our proposed taxonomy is depicted in Figure 1. It characterizes data generation mechanisms using four properties:

1. **Architecture.** Defines the broader type of data augmentation. It is based on domain specificity, architecture type or data transformations using a heuristic or random perturbation process. Generation techniques that apply a form of random perturbation, interpolation or geometric transformation to the data with some degree of randomness are considered randomized approaches. Typical, domain-specific data generation techniques are considered traditional architectures. These techniques apply transformations to a data point using *a priori* domain knowledge. Generative models based on neural network architectures are defined as network-based. These architectures attempt to either generate observations in the feature space and/or by producing observations that are difficult to distinguish from the original dataset.
2. **Application level.** Refers to the phase of the ML pipeline where the generative process is included. Generative models are considered internal if they are used alongside the primary ML task, whereas models used prior to the development of the primary ML task are considered external.
3. **Scope.** Considers the usage of the original dataset’s properties. Generative models that consider the density of the data space, statistical properties of \mathcal{D} , or attempt to replicate specific relationships found in \mathcal{D} are considered to have a global scope, whereas generative models that consider a single observation and/or a set of close neighbors are considered to have a local scope. On the one hand, generative models with a local scope do not account for \mathbb{P}^s but allow for a larger diversity of candidate x^s and higher variance within \mathcal{D}^s . On the other hand, generative models with a global scope have a higher capacity to model \mathbb{P}^s but produce candidate x^s with lower diversity and lower variance within \mathcal{D}^s .

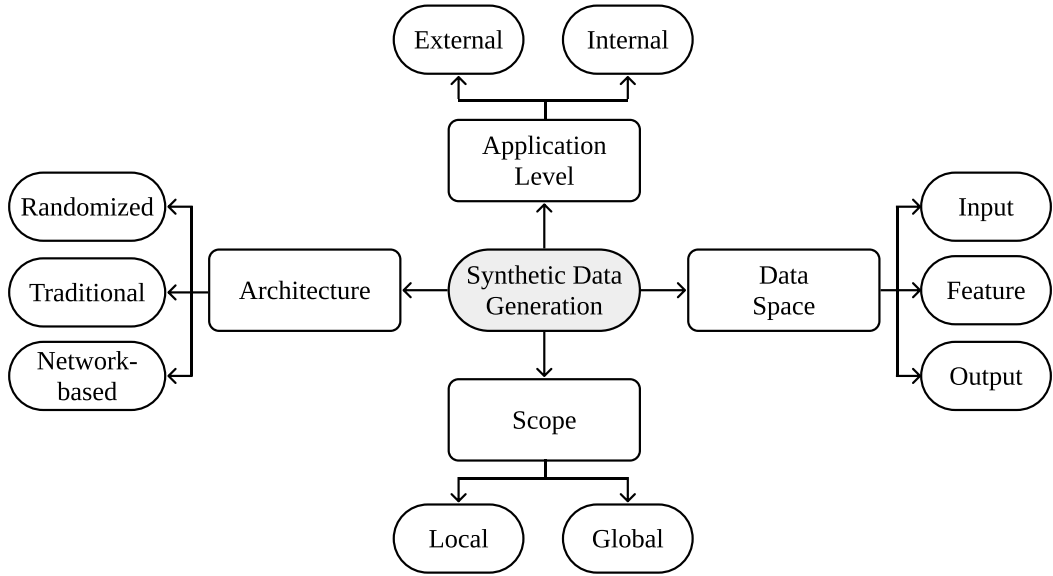


Figure 1: General taxonomy of data generation mechanisms proposed in this paper.

4. Data space. Refers to the type data representation used to apply the generative model. Generation mechanisms can be applied using the raw dataset (*i.e.*, on the input space), an embedded representation of the data (*i.e.*, on the feature space) or based on the target feature (*i.e.*, on the output space). Although some studies discuss the need to generate synthetic data on the input space [30, 31], there are various studies that apply synthetic data generation techniques on a feature space.

Throughout the analysis of the different types of generation mechanisms, all relevant methods were characterized using this taxonomy and listed in Table 2.

Table 2: Summary of the synthetic data generation methods discussed in this work.

Algorithm	ML Problem	Type	Architecture	Level	Data Space	Scope
SynSys [39]	Regression	HMM	Probabilistic	External	Input	Global
SenseGen [40]	Anon. + Reg.	GMM	Net. + Prob.	External	Input	Global
SDV [31]	Anon.	Copula	Probabilistic	External	Input	Global
MST [41]	DP	Marginal	Probabilistic	External	Input	Global
QUAIL [42]	DP	—	—	External	—	Global
SuperQUAIL [43]	DP	—	—	External	—	Global
MWEM [44]	DP	Marginal	Probabilistic	External	Input	Global
MWEM-PGM [45]	DP	PGM	Probabilistic	External	Input	Global
PrivBayes [46]	DP	PGM	Probabilistic	External	Input	Global
DPGAN [47]	DP	GAN	Network	External	Feature	Global
DPCTGAN [42]	DP	GAN	Network	External	Feature	Global
PATE-GAN [48]	DP	GAN	Network	External	Feat. + Out.	Global
PATECTGAN [42]	DP	GAN	Network	External	Feat. + Out.	Global
FEM [49]	DP	Workload	Probabilistic	External	Input	Global
RAP [50]	DP	Workload	Probabilistic	External	Input	Global
PDF [51, 52]	—	—	Probabilistic	External	Input	Global
Kamino [53]	DP	—	Probabilistic	External	Input	Global
RON-GAUSS [54]	DP	Gaussian	Probabilistic	Internal	Feature	Global
HDMM [55]	DP	—	Probabilistic	External	Input	Global

Continued on next page

Table 2: Summary of the synthetic data generation methods discussed in this work.

Algorithm	ML Problem	Type	Architecture	Level	Data Space	Scope
DualQuery [56]	DP		Probabilistic	External	Input	Global
ROS(E) [57]	Ovs	Bootstrap	Randomized	External	Input	Local
SMOTE [58]	Ovs	Linear	Randomized	External	Input	Local
SMOTENC [58]	Ovs					
SMOTEN [58]	Ovs					
Borderline-SMOTE [59]	Ovs	Linear	Randomized	External	Input	Local
G-SMOTE [60]	Ovs	Geometric	Randomized	External	Input	Local
ADASYN [61]	Ovs	Linear	Randomized	External	Input	Local
KernelADASYN [62]	Ovs	Gaussian	Probabilistic	External	Input	Local
MOKAS [63]	Ovs	Rec. Err.	Network	External	Feature	Global
SOMO [64]	Ovs	Linear	Net.+Rand.	External	Input	Global
G-SOMO [65]	Ovs	Geometric	Net.+Rand.	External	Input	Global
Safe-level SMOTE [66]	Ovs	Linear	Randomized	External	Input	Local
LR-SMOTE [67]	Ovs	Linear	Randomized	External	Input	Global
K-means SMOTE [68]	Ovs	Linear	Randomized	External	Input	Global
DBSMOTE [69]	Ovs	Linear	Randomized	External	Input	Local
CGAN [70]	Ovs	GAN	Network	External	Feature	Global
K-means CTGAN [71]	Ovs	GAN	Network	External	Feature	Global
SMOTER [72]	Ovs + Reg					
G-SMOTER [73]	Ovs + Reg					
RACOG [74]	Ovs	PGM	Probabilistic	External	Input	Global
wRACOG [74]	Ovs	PGM	Probabilistic	External	Input	Global
RWO [75]	Ovs	RW	Probabilistic	External	Input	Global
PDFOS [76]	Ovs	PDF	Probabilistic	External	Input	Global
MixUp [77]	DA	Linear			In./Feat.+Out.	
NL-MixUp [78]	DA	Geometric				
AE-DA [79]	DA	AE	Network	External	In./Feat.+Out.	Local
MODALS [80]	DA	—	Network	Internal	Feat.	Global
LSI [81]	DA	AE	Network	External	Feat.+Out.	Global
M-MixUp [82]	DA	AE	Network	External	Feat.+Out.	Global
Gibbs [83]	DA	PGM	Probabilistic	External	Input	Global
MedGAN [84]	DA	GAN	Network	External	Feature	Global
GANBLR [85]	DA	PGM	Probabilistic	External	Input	Global
Table-GAN [86]	DA	GAN	Network	External	Feature	Global
CTGAN [87]	DA	GAN	Network	External	Feature	Global
TVAE [87]	DA	AE	Network	External	Feature	Global
AE [88]	DA	AE	Network	External	Feature	Global
Masking [10]	Self-SL	Mask	Randomized			

4 Generation mechanisms

Laplace perturbations (commonly used as a baseline approach for DP algorithms). Categorical features use n-way marginals (also known as conjunctions or contingency tables [56]) to ensure the generated data contains variability in the categorical features and the distribution of categorical feature values follows some given constraint.

Distribution approximation (discuss marginal inference)

Expectation Maximization algorithm (deterministic)

211 Copula-based mechanisms

212 • Gaussian generative model

213 • Gaussian mixture model

214 • Kernel Density Estimation

215 • Probability density function estimation

216 Probabilistic graphical models

217 • Hidden Markov model

218 • MCMC

219 • Bayesian network

220 • Gibbs sampling

221 Random walk

222 Linear transformations

223 • Linear interpolation

224 – Inter-class interpolation (MixUp)

225 – Inner-class interpolation (SMOTE)

226 • Linear extrapolation

227 • Both

228 Geometric transformations

229 • Hypersphere (G-SMOTE)

230 • Rectangle (NLMixup)

231 Difference transform [\[80\]](#)

232 GANs

233 Autoencoders

234 Reconstruction error-based

235 Random erasing (?)

5 Algorithmic applications

In this section we discuss the data generation mechanisms for the different contexts where they are applied. We emphasize the constraints in each problem that condition the way generation mechanisms are used.

5.1 Privacy

Synthetic data generation is a technique used to produce synthetic, anonymized versions of datasets [30]. It is considered a good approach to share sensitive data without compromising significantly a given data mining task [89, 86]. Traditional data anonymization techniques, as well as federated learning are two other viable solutions for privacy-preserving data publishing tasks, but contain drawbacks [15]. On the one hand, traditional data anonymization requires domain knowledge, is labor intensive and remains susceptible to disclosure [90]. On the other hand, federated learning is a technically complex task that consists on training ML classifiers on edge devices and aggregating temporarily updated parameters on a centralized server, instead of aggregating the training data [91]. Although it prevents sharing sensitive data, its applicability is dependent on the task. Dataset anonymization via synthetic data generation attempts to balance disclosure risk and data utility in the final synthetic dataset. The goal is to ensure observations are not identifiable and the relevant data mining tasks are not compromised [92, 93].

The generation of synthetic datasets allow a more flexible approach to the successful implementation of ML tasks. To do this, it is important to guarantee that sensitive information in \mathcal{D} is not leaked into \mathcal{D}^s . Differential privacy (DP), a formalization of privacy, offers strict theoretical privacy guarantees [42]. A differentially private generation mechanism produces a synthetic dataset, regulated by the privacy parameter ϵ , with statistically indistinguishable results when using either \mathcal{D} or neighboring datasets $\mathcal{D}' = \mathcal{D} \setminus \{x\}$, for any $x \in \mathcal{D}$. A synthetic data generation model (f_{gen}) guarantees (ϵ, δ) -differential privacy if $\forall S \subseteq \text{Range}(f_{gen})$ all $\mathcal{D}, \mathcal{D}'$ differing on a single entry [44]:

$$Pr[f_{gen}(\mathcal{D}) \in S] \leq e^\epsilon \cdot Pr[f_{gen}(\mathcal{D}') \in S] + \delta \quad (1)$$

In this case, ϵ is a non-negative number defined as the privacy budget. A lower ϵ guarantees a higher level of privacy, but reduces the quality of the produced synthetic data. The generation of DP synthetic data is especially appealing since DP is not affected by post-processing; any ML pipeline may be applied using \mathcal{D}^s without losing differential privacy [94].

Despite the formalization and the ability to quantify differential privacy, there are popular synthetic data-based anonymization approaches that perform this task without DP guarantees. Specifically, the Synthetic Data Vault (SDV) [31] is a method for database anonymization that uses Gaussian Copula models for generating data. However, this method allows the usage of other generation mechanisms. A posterior extension of SDV was proposed to generate data using a CTGAN [87] and to handle sequential tabular data using a conditional probabilistic auto-regressive neural network [95].

The choice of the most appropriate DP synthetic data generation techniques depends on the task to be developed (if known) and the domain. However, marginal-based algorithms appear to perform well across various tests [96]. A well-known method for the generation of DP synthetic datasets is the combination of the Multiplicative Weights update rule with the Exponential Mechanism (MWEM) [44]. The MWEM mechanism is an active learning-style algorithm that maintains an approximation of \mathcal{D}^s . At each time step, MWEM selects the worst approximated query (determined by a scoring function) using the Exponential

Mechanism and improves the accuracy of the approximating distribution using the Multiplicative Weights update rule. A known limitation of this method refers to its scalability. Since this method represents the approximate data distribution in datacubes, this method becomes infeasible for high-dimensional problems [45]. This limitation was addressed with the integration of a Probabilistic Graphical Model-based (PGM) estimation into MWEM (MWEM-PGM) and a subroutine to compute and optimize the clique marginals of the PGM, along with other existing privacy mechanisms [45]. Besides MWEM, this method was used to modify and improve the quality of other DP algorithms: PrivBayes [46], HDMM [55] and DualQuery [56].

PrivBayes [46] circumvents the curse of dimensionality by computing a differentially private Bayesian Network (*i.e.*, a type of PGM). Instead of injecting noise into the dataset, they inject noise into the lower-dimensional marginals. The high-dimensional matrix mechanism (HDMM) [55] mechanism is designed to efficiently answer a set of linear queries on high-dimensional data, which are answered using the Laplace mechanism. The DualQuery algorithm [56] is based on the two-player interactions in MWEM, and follows a similar synthetic data generation mechanism as the one found in MWEM.

FEM [49] follows a similar data generation approach as MWEM. It also uses the exponential mechanism and replaces the multiplicative weights update rule with the follow-the-perturbed-leader (FTPL) algorithm [97]. The Relaxed Adaptive Projection (RAP) algorithm [50] uses the projection mechanism [98] to answer queries on the private dataset using a perturbation mechanism and attempts to find the synthetic dataset that matches the noisy answers as accurately as it can.

Kamino [53] introduces denial constraints in the data synthesis process. Kamino builds on top of the probabilistic database framework (PDF) [51, 52], which uses ordinary databases to model a probability distribution and integrates denial constraints as parametric factors, out of which the synthetic observations are sampled. RON-GAUSS [54] combines the random orthonormal (RON) dimensionality reduction technique and synthetic data sampling using either a Gaussian generative model or a Gaussian mixture model. The motivation for this model stems from the *Diaconis-Freedman-Meckes* effect [99], which states that most high-dimensional data projections follow a nearly Gaussian distribution. Since RON-GAUSS includes a feature extraction step (using RON) and the synthetic data generated is not projected back into the input space, we consider RON-GAUSS an internal approach to the ML pipeline.

The MST mechanism [41] is a marginal estimation-based approach that produces differentially private data. It uses the Private-PGM mechanism [45] that relies on the PGM approach to generate synthetic data. PGM models are most commonly used when it is important to maintain the pre-existing statistical properties and relationships between features [100].

The Quail-ified Architecture to Improve Learning (QUAIL) is a DP method that produces differentially private data by distributing the privacy budget between a DP classifier to attribute the target labels onto D^s and the data generator. QUAIL works as a framework that involves the adoption of both a DP classifier and generator. Originally, it was experimented using DPGAN [47], DPCTGAN, MWEM [44], PATE-GAN [48] and PATE-CTGAN. SuperQUAIL [43] is an extension of QUAIL that further distributes the privacy budget according to the feature importance determined using a DP version of SAGE [101]. However, this method does not ensure statistical parity with real data and assumes the task being developed is known *a priori*.

Another family of DP synthetic data generation techniques relies on the usage of Generative Adversarial Networks (GAN). DPGAN [47] modifies the original GAN architecture to make it differentially private by introducing noise to gradients during the learning procedure. This approach was also applied on a conditional GAN architecture directed towards tabular data (CTGAN) [87], which originated the DPCTGAN [42]. Another type of GAN-based DP data synthesis method is based on the combination of

a GAN architecture and the Private Aggregation of Teacher Ensembles (PATE) [102] approach. Although the PATE method generates a DP classifier, it served as the basis for PATE-GAN [48], a DP synthetic data generation mechanism. PATE-GAN replaces the discriminator component of a GAN with the PATE mechanism, which guarantees DP over the generated data. The PATE mechanism is used in the learning phase to train an ensemble of classifiers to distinguish real from synthetic data. In a second step, the predicted labels are passed (with added noise) to another discriminator, which is used to train the generator network.

5.2 Regularization

When there are no underlying problems in the training data, it is sampled from a fixed data source, is labeled, and balanced, the resulting ML classifier is expected to achieve good generalization performance [103]. However, if one or more of these assumptions does not hold, the ML model becomes prone to overfitting [104]. Regularization techniques are often used to address problems like overfitting, small training dataset, high dimensionality, outliers, label noise and catastrophic forgetting [105, 106, 107, 108]. They can be divided into three groups [109]:

1. Output level modifications. Transforms the labels in the training data.
2. Algorithmic level modifications. Modifies the classifier’s architecture, loss function or other components in the training procedure.
3. Input level modifications. Modifies the training dataset by expanding it with synthetic data.

The last approach, input level modifications, is known as data augmentation. It is used to increase the size and data variability of data in a training dataset, by producing synthetic observations [110, 111]. Since it is applied at the data level, it can be used for various types of problems and classifiers [112]. Earlier definitions of data augmentation refer to methods based on iterative optimization or sampling algorithms that introduce unobserved data or latent variables [113]. In the current ML literature, data augmentation techniques mostly refer to the former, while the latter is better known as feature extraction. Although data augmentation is commonly used and extensively studied in computer vision [24] and natural language processing [7], research on tabular data augmentation is sparse.

Feng et al. [79] proposed an autoencoder-based data augmentation (AE-DA) approach where the training of the autoencoder is done for each target class, non-iteratively, which reduces the amount of time required compared to the batch processing approach. The decoding weights of an autoencoder is scaled and linearly combined with an observation from another class using a coefficient that follows the beta distribution. The latter step varies from typical interpolation-based approaches, since this coefficient is usually drawn from a uniform distribution.

The Modality-Agnostic Automated Data Augmentation in the Latent Space model (MODALS) [80] leverages on the concept discussed by DeVries et al. [9], as well as the Latent Space Interpolation method (LSI) [81] and Manifold Mixup [82]. However, MODALS introduces a framework for data augmentation internally. It contains a feature extraction step, trained using a combination of adversarial loss, classification loss and triplet loss, where latent space generation mechanisms are applied. The classifier is trained using the original and the synthetic observations generated in the feature space.

In the model distillation approach proposed in [83] the student model is trained with synthetic data generated with Gibbs sampling. Although Gibbs sampling is infrequently used in recent literature, two

oversampling methods using Gibbs sampling appear to achieve state-of-the-art performance [74]. However, probabilistic-based approaches for data augmentation are uncommon; there are some methods proposed for the more specific case of oversampling, but no more related methods for data augmentation were found.

A well-known approach to GAN-based data augmentation is Table-GAN [86]. It utilizes the vanilla GAN approach to the generation of synthetic data. However, vanilla GAN does not allow the controlled generation of synthetic data given conditional attributes such as the target feature values in supervised learning tasks and may be the cause for aggravated categorical feature imbalance. These limitations were addressed with the CTGAN [87] algorithm, which implements the conditional GAN approach to tabular data. Another GAN-based architecture, MedGAN [84], can also be adapted for tabular data and is used as a benchmark in related studies (*e.g.*, [87, 85]). When compared to the remaining GAN-based approaches, MedGAN’s architecture is more complex and is generally outperformed in the experiments reported in the literature. The GANBLR [85] modifies vanilla GAN architectures with a Bayesian network as both generator and discriminator to create synthetic data that is expected to be indistinguishable from real data. This approach benefits from its interpretability and reduced complexity, while maintaining state-of-the-art performance across various evaluation criteria.

Another less popular approach for network-based synthetic data generation are autoencoder architectures. TVAE, proposed in [87] achieved state-of-the-art performance. It consists of the VAE algorithm with an architecture modified for tabular data (*i.e.*, 1-dimensional). However, as discussed by the authors, this method contains limitations since it is difficult to achieve DP with AE-based models since they access the original data during the training procedure, unlike GANs. Delgado et al. [88] studies the impact of data augmentation on supervised learning with small datasets. The authors compare four different AE architectures: Undercomplete, Sparse, Deep and Variational AE. Although any of the tested AE architectures improved classification performance, the deep and variational autoencoders were the best overall performing models.

Mixup [77]

Automated Data Augmentation — See MODALS paper.

5.3 Time Series

Synsys [39] approaches time-series using both Hidden Markov and regression models. They show the method’s effectiveness in the Healthcare domain with limited ground truth data by comparing it to models trained using only real data. A related model, Sensegen [40], uses an adversarial training approach to train an LSTM that predicts the parameters of Gaussian Mixture Models (GMM) at each time stamp, using real data as an input. Finally, the GMM estimations are used to sample synthetic data.

Generative adversarial networks in time series

Some of the methods previously discussed can also be used for time-series. For example, Cheung et al. [80] show improved performance with time-series data using MixUp and MODALS

One problem frequently found in industry settings is the training of ML models on imbalanced datasets. Since most supervised machine learning classifiers are designed to expect classes with similar frequencies, with highly skewed distributions in \mathcal{D}_L , the classifier’s predictions tend to be biased towards overrepresented classes [3]. For example, one can predict correctly with over 99% accuracy whether credit card accounts were defrauded using a constant classifier. This issue can be addressed in 3 different ways: resampling, algorithmic modifications and cost-sensitive solutions [114]. Resampling techniques are more general approaches when opposed to algorithmic and cost-sensitive methods. They modify \mathcal{D}_L to ensure balanced class frequencies by removing majority class observations (*i.e.*, undersampling), producing synthetic minority class observations (*i.e.*, oversampling), or a combination of both. However, since undersampling removes observations from \mathcal{D}_L , it has the disadvantage of information loss [115] and lacks effectiveness when compared to oversampling methods [116, 117]. Oversampling can be considered a specific setting of data augmentation.

Oversampling is an appropriate technique when, given a set of n target classes, there is a collection C_{maj} containing the majority class observations and C_{min} containing the minority class observations such that $\mathcal{D}_L = \bigcup_{i=1}^n C_i$. The training dataset \mathcal{D}_L is considered imbalanced if $|C_{maj}| > |C_{min}|$. This imbalance is quantified using the Imbalance Ratio (IR), expressed as $IR = \frac{|C_{maj}|}{|C_{min}|}$. An oversampling algorithm with a standard generation policy will generate a $\mathcal{D}_L^s = \bigcup_{i=1}^n C_i^s$ that guarantees $|C_i \cup C_i^s| = |C_{maj}|, \forall i \in \{1, \dots, n\}$. The model f_θ will be trained using an artificially balanced dataset $\mathcal{D}'_L = \mathcal{D}_L \cup \mathcal{D}_L^s$.

Random Oversampling (ROS) is considered a classical approach to oversampling. It oversamples minority classes by randomly picking samples with replacement. It is a bootstrapping approach that, if generated in a smoothed manner (*i.e.*, by adding perturbations to the synthetic data), is also known as Random Oversampling Examples (ROSE) [57]. However, the random duplication of observations often leads to overfitting [118].

The Synthetic Minority Oversampling Technique (SMOTE) [58] attempts to address the data duplication limitation in ROS with a two stage data generation mechanism:

1. Selection phase. A minority class observation, $x^c \in C_{min}$, and one of its k -nearest neighbors, $x^{nn} \in C_{min}$, are randomly selected.
2. Generation phase. A synthetic observation, x^s , is generated along a line segment between x^c and x^{nn} : $x^s = \alpha x^c + (1 - \alpha)x^{nn}, \alpha \sim \mathcal{U}(0, 1)$.

Although the SMOTE algorithm addresses the limitations in ROS, it brings other problems, which motivated the development of several SMOTE-based variants [60]: (1) it introduces noise when a noisy minority class observations is assigned to x^c or x^{nn} , (2) it introduces noise when x^c and x^{nn} belong to different minority-class clusters, (3) it introduces near duplicate observations when x^c and x^{nn} are too close together and (4) it does not account for within-class imbalance (*i.e.*, different input space regions should assume a different importance according to the concentration of minority class observations).

Borderline-SMOTE [59] modifies SMOTE’s selection mechanism. It calculates the k -nearest neighbors for all minority class observations and selects the ones that are going to be used as x^c in the generation phase. An observation is selected based on the number of neighbors belonging to a different class, where the observations with no neighbors belonging to C_{min} and insufficient number of neighbors belonging to C_{maj} are not considered for the generation phase. This approximates the synthetic observations to the

border of the expected decision boundaries. Various other methods were proposed since then to modify selection mechanism, such as K-means SMOTE [68]. This approach addresses within-class imbalance and the generation of noisy synthetic data by generating data within clusters. The data generation is done according to each cluster’s imbalance ratio and dispersion of minority class observations. DBSMOTE [69] also modifies the selection strategy by selecting as x^c the set of core observations in a DBSCAN clustering solution.

The Adaptive Synthetic Sampling approach (ADASYN) [61] uses a comparable approach to Borderline-SMOTE. It calculates the ratio of non-minority class observations within the k -nearest neighbors of each $x \in C_{min}$. The amount of observations to be generated using each $x \in C_{min}$ as x^c is determined according to this ratio; the more non-minority class neighbors an observation contains, the more synthetic observations are generated using it as x^c . The generation phase is done using the linear mechanism in SMOTE. However, this approach tends to aggravate the limitation (1) previously discussed. A second version of this method, KernelADASYN [62], replaces the generation mechanism with a weighted kernel density estimation. The weighing is done according to ADASYN’s ratio and the synthetic data is sampled using the calculated Gaussian Kernel function whose bandwidth is passed as an additional hyperparameter.

Modifications to SMOTE’s generation mechanism are less common and generally attempt to address problem of noisy synthetic data generation. Safe-level SMOTE [66] truncates the line segment between x^c and x^{nn} according to a safe level ratio. Geometric-SMOTE (G-SMOTE) [60] it generates synthetic data within a deformed and truncated hypersphere to also avoid the generation of near-duplicate synthetic data. It also introduces a modification of the selection strategy to combine the selection of majority class observations as x^{nn} to avoid the introduction of noisy synthetic data.

LR-SMOTE [67] modifies both the selection and generation mechanisms. The set of observations to use as x^c contains the misclassified minority class observations using a SVM classifier, out of which the potentially noisy observations are removed. The k-means clustering method is used to find the closest observations to the cluster centroids, which are used as x^c . The observations with a higher number of majority class neighbors are more likely to be selected as x^{nn} . Although the generation mechanism synthesizes observations as a linear combination between x^c and x^{nn} , it restricts or expands this range by setting $\alpha \sim \mathcal{U}(0, M)$, where M is a ratio between the average euclidean distance of each cluster’s minority class observations to x^c and the euclidean distance between x^c and x^{nn} .

The Minority Oversampling Kernel Adaptive Subspaces algorithm (MOKAS) [63] adopts a different approach when compared to SMOTE-based mechanisms. It uses the adaptive subspace self-organizing map (ASSOM) [119] algorithm to learn sub-spaces (*i.e.*, different feature spaces for each unit in the SOM), out of which synthetic data is generated. The synthetic data is generated using a lower dimensional representation of the input data to ensure the reconstructed data is different from the original observations. Overall, the usage of SOMs for oversampling is uncommon. Another two examples of this approach, SOMO [64] and G-SOMO [65] use a similar approach as K-means SMOTE. In the case of G-SOMO, instead of using SMOTE’s generation mechanism, it uses G-SMOTE’s instead.

Another set of network-based methods that fully replace SMOTE-based mechanisms are GAN-based architectures. One example of this approach is CGAN [70]. It uses an adversarial training approach to generate data that approximates the original data distribution and indistinguishable from the original dataset (according to the adversarial classifier). A more recent GAN-based oversampler, K-means CTGAN [71] uses a K-means clustering method as an additional attribute to train the CTGAN. In this case, cluster labels allow the reduction of within-class imbalance. These types of approaches benefit from learning the overall per-class distribution, instead of using local information only. However, GANs require more computational power to train, their performance is sensitive to the initialization and are prone to

482 the “mode collapse” problem.

483 Statistical-based oversampling approaches are less common. Some methods, such as RACOG and
484 wRACOG [74] are based on Gibbs sampling, PDFOS [76] is based on probability density function
485 estimations and RWO [75] uses a random walk algorithm.

486 5.5 Active Learning

487

488 5.6 Few-shot Learning

489

490 Analysis of six feature space data augmentation techniques for few-shot learning [37]

491 FlipDA [120]

492 Data generation can be used to address Few-shot learning in three ways [121]: (1) transforming samples
493 from the dataset, (2) transforming samples from a weakly labeled or unlabeled dataset, or (3) transforming
494 samples from similar datasets.

495 5.7 Semi-supervised Learning

496

497 Synthetic data generation for semi-supervised learning given limited labeled data regarding the COVID-19
498 pandemic [122].

499 Extensive literature review on semi-supervised learning [38]

500 5.8 Self-supervised Learning

501

502 Masking [10]

503 SubTab [123]

504 6 Evaluating the Quality of Synthetic Data

505

506 The log-likelihood (and equivalently the Kullback-Leibler Divergence) is a de-facto standard to train and
507 evaluate generative models [12]. Other common metrics include Parzen window estimates, which Theis
508 et al. [12] show that these metrics behave independently and should generally be avoided. Therefore, it is
509 necessary to evaluate generative models with respect to the application these models are being developed
510 for.

511 The evaluation of generative models should quantify three key aspects of synthetic data [13]:

- 512 1. Fidelity
- 513 2. Diversity
- 514 3. Generalization

515 The 3-dimensional metric proposed by Alaa et al. [13] quantifies these aspects via the combination of
516 three metrics (α -Precision, β -Recall and Authenticity) for various application domains.

517 The vast majority of synthetic data generation models are evaluated on a ML utility basis. However, the
518 GANBLR model [85] was evaluated based on three aspects: (1) ML utility, (2) Statistical similarity, and
519 (3) Interpretability. In Xu et al. [87], the authors evaluate the proposed method using a likelihood fitness
520 metric (to measure statistical similarity) and ML efficacy (*i.e.*, utility).

521 6.1 Statistical Divergence Metrics

522 6.2 Precision/Recall Metrics

523 6.3 Supervised Learning Metrics

524 7 Discussion

525

526 7.1 Main Findings

527 The combination of data generation strategies is an approach commonly found in different problems, such
528 as self-supervised learning [5]. It can be more frequently found in text data applications [18] and image
529 data [CITATION].

530 **7.1.1 RQ1: bla bla bla**

531 **7.1.2 RQ2: bla bla bla**

532 **7.1.3 RQ3: bla bla bla**

533 **7.2 Limitations**

534 Research across the different applications appears to be sandboxed even though all techniques integrate
535 synthetic data in its core.

536 It is generally understood that, if learned properly, the feature space is expected to be convex and isotropic.
537 In that case, using linear generation techniques in the feature space would produce synthetic data without
538 introducing noise [80]. However, it is unclear which types of model/architectures and training procedures
539 contribute to the learning of a good feature space according to the context.

540 Given the breadth and complexity of input-level and feature-level data generation mechanisms, it is
541 increasingly important to find a method to efficiently determine the most appropriate data generation
542 policies. However, the complexity of this task is determined by various factors: different data types, ML
543 problems, model architectures, computational resources, performance metrics and contextual constraints.
544 Auto-augmentation and meta learning aim to address this challenge and are still subject to active
545 research.

546 The quality of synthetic data generation in high-dimensional domains appears as a prevailing limitation
547 in most applications. This method might be addressed with dimensionality reduction techniques along
548 with data generation in the feature space. However, research on generation in the feature space is greatly
549 focused on GAN architectures, which require significant computational power. Other methods for learning
550 manifold space embeddings could be explored to address this limitation.

551 To the best of our knowledge, research on data augmentation using auto-encoder architectures is sparse.
552 There is, however, a few studies performing data augmentation in different domains using tabular data [88].
553 More commonly, autoencoders are used to learn a manifold features space for more complex data types.
554 As long as the method used to generate the feature space is appropriate, the methods discussed in this
555 study can be used in the feature space regardless of the type of data.

556 It remains an open question which feature space transformations, or types of transformations, create
557 better synthetic data [80].

558 There is not much research concerning the quality and general performance between data generation on
559 the input, feature and output space.

560 The evaluation of anonymization techniques lack standardized, objective and reliable performance metrics
561 and benchmark datasets to allow an easier comparison across classifiers to evaluate key aspects of data
562 anonymization (resemblance, utility, privacy and performance). These datasets should contain mixed data
563 types (*i.e.*, a combination of categorical, ordinal, continuous and discrete features) and the metrics should
564 evaluate the performance of different data mining tasks along with the anonymization reliability. This
565 problem appears to be universal across domains. For example, Hernandez et al. [15] observed the lack of
566 a universal method or metric to report the performance synthetic data generation algorithms for tabular

567 health records. Therefore, in order to facilitate the usage of these techniques in industry domains, these
568 benchmarks must also be realistic. Rosenblatt et al. [42] attempts to address this problem by proposing a
569 standardized evaluation methodology using standard datasets and real-world industry applications.

570 Computational cost and inconsistent quality of synthetic data generated with GANs (*e.g.*, mode collapse).

571 Research on differentially private variational autoencoders is sparse to non-existent. The only related
572 study found in the literature was developed in [124]. However, it is not peer reviewed or particularly
573 popular, which led us to discard this paper from our analysis.

574 Unlike with data privacy solutions, data augmentation techniques generally do not consider the simi-
575 larity/dissimilarity of synthetic data. The study of quality metrics for supervised learning may reduce
576 computational overhead and experimentation time. No studies related to the relationship of quality
577 metrics and performance in the primary ML task were found [CONFIRM!!!].

578 There is not a clear understanding of what types of data augmentation methods are more appropriate
579 according to different model architectures, ML tasks or domains and the reason why they work better or
580 worse depending on the task. In addition, it is still unclear *why* data augmentation works. Research on
581 this topic lacks depth and fails to address the theoretical underpinnings [7].

582 In some domains, a common approach for data augmentation is the combination of several data augmenta-
583 tion methods to increase the diversification of synthetic data. This is true for both text classification [18]
584 and image classification [CITATION]. However, for tabular data, no similar approach was found.

585 “Dao et al. (2019) note that “data augmentation is typically performed in an ad-hoc manner with little
586 understanding of the underlying theoretical principles”, and claim the typical explanation of DA as
587 regularization to be insufficient.” [7]

588 There is a lack of research on oversampling solutions to generate synthetic data with mixed data types
589 and datasets with exclusively non metric features.

590 There is a lack of methods adapted to use categorical features for tabular data.

591 There is a lack of methods directed to regression problems.

592 There is a paucity of research on the usage of probabilistic-based generation mechanisms in oversampling.

593 There is no clear understanding of the most appropriate data augmentation techniques used to train
594 self-supervised models and how their behavior and performance varies according to the data generation
595 method used.

596 Oversampling does not seem to be a relevant source of bias in behavioral research and does not appear to
597 have an appreciably different effect on results for directly versus indirectly oversampled variables [125].
598 However, most oversampling methods do not account for the distribution in \mathcal{D} , which is especially
599 important for features with sensitive information (*e.g.*, gender or ethnicity). Therefore, the application of
600 oversampling methods on user data may further increase the bias in classification/discrimination between
601 gender or ethnicity groups.

7.3 Research directions

Quantifying the quality of the generated data:

1. Realistic
2. Similarity
3. Usefulness (determine purpose and relevant performance metric)
4. Understand the relationship between the 3 factors

8 Conclusions

References

- [1] Samuel A Assefa, Danial Dervovic, Mahmoud Mahfouz, Robert E Tillman, Prashant Reddy, and Manuela Veloso. “Generating synthetic data in finance: opportunities, challenges and pitfalls”. In: *Proceedings of the First ACM International Conference on AI in Finance*. 2020, pp. 1–8.
- [2] Samuli Laine and Timo Aila. “Temporal ensembling for semi-supervised learning”. In: *International Conference on Learning Representations (ICLR)*. Vol. 4. 5. 2017, p. 6.
- [3] Joao Fonseca, Georgios Douzas, and Fernando Bacao. “Improving imbalanced land cover classification with K-Means SMOTE: Detecting and oversampling distinctive minority spectral signatures”. In: *Information* 12.7 (2021), p. 266.
- [4] Yoon-Yeong Kim, Kyungwoo Song, JoonHo Jang, and Il-Chul Moon. “LADA: Look-Ahead Data Acquisition via Augmentation for Deep Active Learning”. In: *Advances in Neural Information Processing Systems* 34 (2021), pp. 22919–22930.
- [5] Jean-Bastien Grill, Florian Strub, Florent Altché, Corentin Tallec, Pierre Richemond, Elena Buchatskaya, Carl Doersch, Bernardo Avila Pires, Zhaohan Guo, Mohammad Gheshlaghi Azar, et al. “Bootstrap your own latent-a new approach to self-supervised learning”. In: *Advances in neural information processing systems* 33 (2020), pp. 21271–21284.
- [6] Jiang-Jing Lv, Xiao-Hu Shao, Jia-Shui Huang, Xiang-Dong Zhou, and Xi Zhou. “Data augmentation for face recognition”. In: *Neurocomputing* 230 (2017), pp. 184–196.
- [7] Steven Y Feng, Varun Gangal, Jason Wei, Sarath Chandar, Soroush Vosoughi, Teruko Mitamura, and Eduard Hovy. “A Survey of Data Augmentation Approaches for NLP”. In: *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*. Online: Association for Computational Linguistics, Aug. 2021, pp. 968–988. DOI: [10.18653/v1/2021.findings-acl.84](https://doi.org/10.18653/v1/2021.findings-acl.84). URL: <https://aclanthology.org/2021.findings-acl.84>.
- [8] Talha Mahboob Alam, Kamran Shaukat, Ibrahim A Hameed, Suhui Luo, Muhammad Umer Sarwar, Shakir Shabbir, Jiaming Li, and Matloob Khushi. “An investigation of credit card default prediction in the imbalanced datasets”. In: *IEEE Access* 8 (2020), pp. 201173–201198.

- [9] Terrance DeVries and Graham W Taylor. “Dataset augmentation in feature space”. In: *arXiv preprint arXiv:1702.05538* (2017).
- [10] Jinsung Yoon, Yao Zhang, James Jordon, and Mihaela van der Schaar. “Vime: Extending the success of self-and semi-supervised learning to tabular domain”. In: *Advances in Neural Information Processing Systems* 33 (2020), pp. 11033–11043.
- [11] Diederik P Kingma, Max Welling, et al. “An introduction to variational autoencoders”. In: *Foundations and Trends® in Machine Learning* 12.4 (2019), pp. 307–392.
- [12] L Theis, A van den Oord, and M Bethge. “A note on the evaluation of generative models”. In: *International Conference on Learning Representations (ICLR 2016)*. 2016, pp. 1–10.
- [13] Ahmed Alaa, Boris Van Breugel, Evgeny S Saveliev, and Mihaela van der Schaar. “How faithful is your synthetic data? sample-level metrics for evaluating and auditing generative models”. In: *International Conference on Machine Learning*. PMLR. 2022, pp. 290–306.
- [14] Miro Mannino and Azza Abouzied. “Is this real? Generating synthetic data that looks real”. In: *Proceedings of the 32nd Annual ACM Symposium on User Interface Software and Technology*. 2019, pp. 549–561.
- [15] Mikel Hernandez, Gorka Epelde, Ane Alberdi, Rodrigo Cilla, and Debbie Rankin. “Synthetic Data Generation for Tabular Health Records: A Systematic Review”. In: *Neurocomputing* (2022).
- [16] Trivellore E Raghunathan. “Synthetic data”. In: *Annual Review of Statistics and Its Application* 8 (2021), pp. 129–140.
- [17] Jakub Nalepa, Michal Marcinkiewicz, and Michal Kawulok. “Data augmentation for brain-tumor segmentation: a review”. In: *Frontiers in computational neuroscience* 13 (2019), p. 83.
- [18] Markus Bayer, Marc-André Kaufhold, and Christian Reuter. “A survey on data augmentation for text classification”. In: *ACM Computing Surveys* (2021).
- [19] Connor Shorten, Taghi M Khoshgoftaar, and Borko Furht. “Text data augmentation for deep learning”. In: *Journal of big Data* 8.1 (2021), pp. 1–34.
- [20] Jiaao Chen, Derek Tam, Colin Raffel, Mohit Bansal, and Diyi Yang. “An empirical survey of data augmentation for limited data learning in NLP”. In: *arXiv preprint arXiv:2106.07499* (2021).
- [21] Pei Liu, Xuemin Wang, Chao Xiang, and Weiye Meng. “A survey of text data augmentation”. In: *2020 International Conference on Computer Communication and Network Security (CCNS)*. IEEE. 2020, pp. 191–195.
- [22] Xin Yi, Ekta Walia, and Paul Babyn. “Generative adversarial network in medical imaging: A review”. In: *Medical image analysis* 58 (2019), p. 101552.
- [23] Xiang Wang, Kai Wang, and Shiguo Lian. “A survey on face data augmentation for the training of deep neural networks”. In: *Neural computing and applications* 32.19 (2020), pp. 15503–15531.
- [24] Connor Shorten and Taghi M Khoshgoftaar. “A survey on image data augmentation for deep learning”. In: *Journal of big data* 6.1 (2019), pp. 1–48.
- [25] Cherry Khosla and Baljit Singh Saini. “Enhancing performance of deep learning models with different data augmentation techniques: A survey”. In: *2020 International Conference on Intelligent Engineering and Management (ICIEM)*. IEEE. 2020, pp. 79–85.
- [26] Brian Kenji Iwana and Seiichi Uchida. “An empirical survey of data augmentation for time series classification with neural networks”. In: *Plos one* 16.7 (2021), e0254841.
- [27] Qingsong Wen, Liang Sun, Fan Yang, Xiaomin Song, Jingkun Gao, Xue Wang, and Huan Xu. “Time series data augmentation for deep learning: a survey”. In: *Proceedings of the Thirtieth International Joint Conference on Artificial Intelligence, IJCAI-21*. Ed. by Zhi-Hua Zhou. International Joint Conferences on Artificial Intelligence Organization, Aug. 2021, pp. 4653–4660.

- [28] Tong Zhao, Gang Liu, Stephan Günnemann, and Meng Jiang. “Graph Data Augmentation for Graph Machine Learning: A Survey”. In: *arXiv preprint arXiv:2202.08871* (2022).
- [29] Nour Eldeen Khalifa, Mohamed Loey, and Seyedali Mirjalili. “A comprehensive survey of recent trends in deep learning for digital images augmentation”. In: *Artificial Intelligence Review* (2021), pp. 1–27.
- [30] Fida K Dankar and Mahmoud Ibrahim. “Fake it till you make it: Guidelines for effective synthetic data generation”. In: *Applied Sciences* 11.5 (2021), p. 2158.
- [31] Neha Patki, Roy Wedge, and Kalyan Veeramachaneni. “The synthetic data vault”. In: *2016 IEEE International Conference on Data Science and Advanced Analytics (DSAA)*. IEEE. 2016, pp. 399–410.
- [32] Chiyuan Zhang, Samy Bengio, Moritz Hardt, Benjamin Recht, and Oriol Vinyals. “Understanding deep learning (still) requires rethinking generalization”. In: *Communications of the ACM* 64.3 (2021), pp. 107–115.
- [33] Yi Zeng, Han Qiu, Gerard Memmi, and Meikang Qiu. “A data augmentation-based defense method against adversarial attacks in neural networks”. In: *International Conference on Algorithms and Architectures for Parallel Processing*. Springer. 2020, pp. 274–289.
- [34] John X Morris, Eli Lifland, Jin Yong Yoo, Jake Grigsby, Di Jin, and Yanjun Qi. “Textattack: A framework for adversarial attacks, data augmentation, and adversarial training in nlp”. In: *arXiv preprint arXiv:2005.05909* (2020).
- [35] José A Sáez, Bartosz Krawczyk, and Michał Woźniak. “Analyzing the oversampling of different classes and types of examples in multi-class imbalanced datasets”. In: *Pattern Recognition* 57 (2016), pp. 164–178.
- [36] Joao Fonseca, Georgios Douzas, and Fernando Bacao. “Increasing the Effectiveness of Active Learning: Introducing Artificial Data Generation in Active Learning for Land Use/Land Cover Classification”. In: *Remote Sensing* 13.13 (2021), p. 2619.
- [37] Varun Kumar, Hadrien Glaude, Cyprien de Lichy, and William Campbell. “A Closer Look At Feature Space Data Augmentation For Few-Shot Intent Classification”. In: *Proceedings of the 2nd Workshop on Deep Learning Approaches for Low-Resource NLP (DeepLo 2019)*. 2019, pp. 1–10.
- [38] Jesper E Van Engelen and Holger H Hoos. “A survey on semi-supervised learning”. In: *Machine Learning* 109.2 (2020), pp. 373–440.
- [39] Jessamyn Dahmen and Diane Cook. “SynSys: A synthetic data generation system for healthcare applications”. In: *Sensors* 19.5 (2019), p. 1181.
- [40] Moustafa Alzantot, Supriyo Chakraborty, and Mani Srivastava. “Sensegen: A deep learning architecture for synthetic sensor data generation”. In: *2017 IEEE International Conference on Pervasive Computing and Communications Workshops (PerCom Workshops)*. IEEE. 2017, pp. 188–193.
- [41] Ryan McKenna, Gerome Miklau, and Daniel Sheldon. “Winning the NIST Contest: A scalable and general approach to differentially private synthetic data”. In: *Journal of Privacy and Confidentiality* 11.3 (2021).
- [42] Lucas Rosenblatt, Xiaoyan Liu, Samira Pouyanfar, Eduardo de Leon, Anuj Desai, and Joshua Allen. “Differentially private synthetic data: Applied evaluations and enhancements”. In: *arXiv preprint arXiv:2011.05537* (2020).
- [43] Lucas Rosenblatt, Joshua Allen, and Julia Stoyanovich. “Spending Privacy Budget Fairly and Wisely”. In: *arXiv preprint arXiv:2204.12903* (2022).
- [44] Moritz Hardt, Katrina Ligett, and Frank McSherry. “A simple and practical algorithm for differentially private data release”. In: *Proceedings of the 25th International Conference on Neural Information Processing Systems-Volume 2*. 2012, pp. 2339–2347.

- [45] Ryan McKenna, Daniel Sheldon, and Gerome Miklau. “Graphical-model based estimation and inference for differential privacy”. In: *International Conference on Machine Learning*. PMLR. 2019, pp. 4435–4444.
- [46] Jun Zhang, Graham Cormode, Cecilia M Procopiuc, Divesh Srivastava, and Xiaokui Xiao. “Privbayes: Private data release via bayesian networks”. In: *ACM Transactions on Database Systems (TODS)* 42.4 (2017), pp. 1–41.
- [47] Liyang Xie, Kaixiang Lin, Shu Wang, Fei Wang, and Jiayu Zhou. “Differentially private generative adversarial network”. In: *arXiv preprint arXiv:1802.06739* (2018).
- [48] James Jordon, Jinsung Yoon, and Mihaela Van Der Schaar. “PATE-GAN: Generating synthetic data with differential privacy guarantees”. In: *International conference on learning representations*. 2018.
- [49] Giuseppe Vietri, Grace Tian, Mark Bun, Thomas Steinke, and Steven Wu. “New oracle-efficient algorithms for private synthetic data release”. In: *International Conference on Machine Learning*. PMLR. 2020, pp. 9765–9774.
- [50] Sergul Aydoore, William Brown, Michael Kearns, Krishnaram Kenthapadi, Luca Melis, Aaron Roth, and Ankit A Siva. “Differentially private query release through adaptive projection”. In: *International Conference on Machine Learning*. PMLR. 2021, pp. 457–467.
- [51] Christopher De Sa, Ihab Ilyas, Benny Kimelfeld, Christopher Re, and Theodoros Rekatsinas. “A Formal Framework for Probabilistic Unclean Databases”. In: *22nd International Conference on Database Theory (ICDT 2019)*. 2019.
- [52] Dan Suciu, Dan Olteanu, Christopher Ré, and Christoph Koch. “Probabilistic databases”. In: *Synthesis lectures on data management* 3.2 (2011), pp. 1–180.
- [53] Chang Ge, Shubhankar Mohapatra, Xi He, and Ihab F Ilyas. “Kamino: constraint-aware differentially private data synthesis”. In: *Proceedings of the VLDB Endowment* 14.10 (2021), pp. 1886–1899.
- [54] Thee Chanyaswad, Changchang Liu, and Prateek Mittal. “Ron-gauss: Enhancing utility in non-interactive private data release”. In: *Proceedings on Privacy Enhancing Technologies* 2019.1 (2019), pp. 26–46.
- [55] Ryan McKenna, Gerome Miklau, Michael Hay, and Ashwin Machanavajjhala. “Optimizing error of high-dimensional statistical queries under differential privacy”. In: *Proceedings of the VLDB Endowment* 11.10 (2018).
- [56] Marco Gaboardi, Emilio Jesús Gallego Arias, Justin Hsu, Aaron Roth, and Zhiwei Steven Wu. “Dual query: Practical private query release for high dimensional data”. In: *International Conference on Machine Learning*. PMLR. 2014, pp. 1170–1178.
- [57] Giovanna Menardi and Nicola Torelli. “Training and assessing classification rules with imbalanced data”. In: *Data mining and knowledge discovery* 28.1 (2014), pp. 92–122.
- [58] Nitesh V Chawla, Kevin W Bowyer, Lawrence O Hall, and W Philip Kegelmeyer. “SMOTE: synthetic minority over-sampling technique”. In: *Journal of artificial intelligence research* 16 (2002), pp. 321–357.
- [59] Hui Han, Wen-Yuan Wang, and Bing-Huan Mao. “Borderline-SMOTE: a new over-sampling method in imbalanced data sets learning”. In: *International conference on intelligent computing*. Springer. 2005, pp. 878–887.
- [60] Georgios Douzas and Fernando Bacao. “Geometric SMOTE a geometrically enhanced drop-in replacement for SMOTE”. In: *Information Sciences* 501 (2019), pp. 118–135.
- [61] Haibo He, Yang Bai, Eduardo A Garcia, and Shutao Li. “ADASYN: Adaptive synthetic sampling approach for imbalanced learning”. In: *2008 IEEE international joint conference on neural networks (IEEE world congress on computational intelligence)*. IEEE. 2008, pp. 1322–1328.

- [62] Bo Tang and Haibo He. “KernelADASYN: Kernel based adaptive synthetic data generation for imbalanced learning”. In: *2015 IEEE congress on evolutionary computation (CEC)*. IEEE. 2015, pp. 664–671.
- [63] Chin-Teng Lin, Tsung-Yu Hsieh, Yu-Ting Liu, Yang-Yin Lin, Chieh-Ning Fang, Yu-Kai Wang, Gary Yen, Nikhil R Pal, and Chun-Hsiang Chuang. “Minority oversampling in kernel adaptive subspaces for class imbalanced datasets”. In: *IEEE Transactions on Knowledge and Data Engineering* 30.5 (2017), pp. 950–962.
- [64] Georgios Douzas and Fernando Bacao. “Self-Organizing Map Oversampling (SOMO) for imbalanced data set learning”. In: *Expert systems with Applications* 82 (2017), pp. 40–52.
- [65] Georgios Douzas, Rene Rauch, and Fernando Bacao. “G-SOMO: An oversampling approach based on self-organized maps and geometric SMOTE”. In: *Expert Systems with Applications* 183 (2021), p. 115230.
- [66] Chumphol Bunkhumpornpat, Krung Sinapiromsaran, and Chidchanok Lursinsap. “Safe-level-smote: Safe-level-synthetic minority over-sampling technique for handling the class imbalanced problem”. In: *Pacific-Asia conference on knowledge discovery and data mining*. Springer. 2009, pp. 475–482.
- [67] XW Liang, AP Jiang, T Li, YY Xue, and GT Wang. “LR-SMOTE—An improved unbalanced data set oversampling based on K-means and SVM”. In: *Knowledge-Based Systems* 196 (2020), p. 105845.
- [68] Georgios Douzas, Fernando Bacao, and Felix Last. “Improving imbalanced learning through a heuristic oversampling method based on k-means and SMOTE”. In: *Information Sciences* 465 (2018), pp. 1–20.
- [69] Chumphol Bunkhumpornpat, Krung Sinapiromsaran, and Chidchanok Lursinsap. “DBSMOTE: density-based synthetic minority over-sampling technique”. In: *Applied Intelligence* 36.3 (2012), pp. 664–684.
- [70] Georgios Douzas and Fernando Bacao. “Effective data generation for imbalanced learning using conditional generative adversarial networks”. In: *Expert Systems with applications* 91 (2018), pp. 464–471.
- [71] Chunsheng An, Jingtong Sun, Yifeng Wang, and Qingjie Wei. “A K-means Improved CTGAN Oversampling Method for Data Imbalance Problem”. In: *2021 IEEE 21st International Conference on Software Quality, Reliability and Security (QRS)*. IEEE. 2021, pp. 883–887.
- [72] Luís Torgo, Rita P Ribeiro, Bernhard Pfahringer, and Paula Branco. “Smote for regression”. In: *Portuguese conference on artificial intelligence*. Springer. 2013, pp. 378–389.
- [73] Luís Camacho, Georgios Douzas, and Fernando Bacao. “Geometric SMOTE for regression”. In: *Expert Systems with Applications* (2022), p. 116387.
- [74] Barnan Das, Narayanan C Krishnan, and Diane J Cook. “RACOG and wRACOG: Two probabilistic oversampling techniques”. In: *IEEE transactions on knowledge and data engineering* 27.1 (2014), pp. 222–234.
- [75] Huaxiang Zhang and Mingfang Li. “RWO-Sampling: A random walk over-sampling approach to imbalanced data classification”. In: *Information Fusion* 20 (2014), pp. 99–116.
- [76] Ming Gao, Xia Hong, Sheng Chen, Chris J Harris, and Emad Khalaf. “PDFOS: PDF estimation based over-sampling for imbalanced two-class problems”. In: *Neurocomputing* 138 (2014), pp. 248–259.
- [77] Hongyi Zhang, Moustapha Cisse, Yann N Dauphin, and David Lopez-Paz. “mixup: Beyond Empirical Risk Minimization”. In: *International Conference on Learning Representations*. 2018.
- [78] Hongyu Guo. “Nonlinear mixup: Out-of-manifold data augmentation for text classification”. In: *Proceedings of the AAAI Conference on Artificial Intelligence*. Vol. 34. 04. 2020, pp. 4044–4051.

- [79] Xiexing Feng, QM Jonathan Wu, Yimin Yang, and Libo Cao. “An autoencoder-based data augmentation strategy for generalization improvement of DCNNs”. In: *Neurocomputing* 402 (2020), pp. 283–297.
- [80] Tsz-Him Cheung and Dit-Yan Yeung. “Modals: Modality-agnostic automated data augmentation in the latent space”. In: *International Conference on Learning Representations*. 2020.
- [81] Xiaofeng Liu, Yang Zou, Lingsheng Kong, Zhihui Diao, Junliang Yan, Jun Wang, Site Li, Ping Jia, and Jane You. “Data augmentation via latent space interpolation for image classification”. In: *2018 24th International Conference on Pattern Recognition (ICPR)*. IEEE. 2018, pp. 728–733.
- [82] Vikas Verma, Alex Lamb, Christopher Beckham, Amir Najafi, Ioannis Mitliagkas, David Lopez-Paz, and Yoshua Bengio. “Manifold mixup: Better representations by interpolating hidden states”. In: *International Conference on Machine Learning*. PMLR. 2019, pp. 6438–6447.
- [83] Rasool Fakoor, Jonas W Mueller, Nick Erickson, Pratik Chaudhari, and Alexander J Smola. “Fast, accurate, and simple models for tabular data via augmented distillation”. In: *Advances in Neural Information Processing Systems* 33 (2020), pp. 8671–8681.
- [84] Karim Armanious, Chenming Jiang, Marc Fischer, Thomas Küstner, Tobias Hepp, Konstantin Nikolaou, Sergios Gatidis, and Bin Yang. “MedGAN: Medical image translation using GANs”. In: *Computerized medical imaging and graphics* 79 (2020), p. 101684.
- [85] Yishuo Zhang, Nayyar A Zaidi, Jiahui Zhou, and Gang Li. “GANBLR: a tabular data generation model”. In: *2021 IEEE International Conference on Data Mining (ICDM)*. IEEE. 2021, pp. 181–190.
- [86] Noseong Park, Mahmoud Mohammadi, Kshitij Gorde, Sushil Jajodia, Hongkyu Park, and Youngmin Kim. “Data Synthesis based on Generative Adversarial Networks”. In: *Proceedings of the VLDB Endowment* 11.10 (2018).
- [87] Lei Xu, Maria Skoularidou, Alfredo Cuesta-Infante, and Kalyan Veeramachaneni. “Modeling tabular data using conditional gan”. In: *Advances in Neural Information Processing Systems* 32 (2019).
- [88] Juan Manuel Davila Delgado and Lukumon Oyedele. “Deep learning with small datasets: using autoencoders to address limited datasets in construction management”. In: *Applied Soft Computing* 112 (2021), p. 107836.
- [89] Jennifer Taub, Mark Elliot, Maria Pampaka, and Duncan Smith. “Differential correct attribution probability for synthetic data: an exploration”. In: *International Conference on Privacy in Statistical Databases*. Springer. 2018, pp. 122–137.
- [90] Jerome P Reiter. “New approaches to data dissemination: A glimpse into the future (?)” In: *Chance* 17.3 (2004), pp. 11–15.
- [91] Bin Yu, Wenjie Mao, Yihan Lv, Chen Zhang, and Yu Xie. “A survey on federated learning in data mining”. In: *Wiley Interdisciplinary Reviews: Data Mining and Knowledge Discovery* 12.1 (2022), e1443.
- [92] Kalpana Singh and Lynn Batten. “Aggregating privatized medical data for secure querying applications”. In: *Future Generation Computer Systems* 72 (2017), pp. 250–263.
- [93] Ping Li, Tong Li, Heng Ye, Jin Li, Xiaofeng Chen, and Yang Xiang. “Privacy-preserving machine learning with multiple data providers”. In: *Future Generation Computer Systems* 87 (2018), pp. 341–350.
- [94] Cynthia Dwork, Aaron Roth, et al. “The algorithmic foundations of differential privacy”. In: *Foundations and Trends® in Theoretical Computer Science* 9.3–4 (2014), pp. 211–407.
- [95] Kevin Zhang, Neha Patki, and Kalyan Veeramachaneni. “Sequential Models in the Synthetic Data Vault”. In: *arXiv preprint arXiv:2207.14406* (2022).

- [96] Yuchao Tao, Ryan McKenna, Michael Hay, Ashwin Machanavajjhala, and Gerome Miklau. “Benchmarking differentially private synthetic data generation algorithms”. In: *arXiv e-prints* (2021), arXiv-2112.
- [97] Adam Kalai and Santosh Vempala. “Efficient algorithms for online decision problems”. In: *Journal of Computer and System Sciences* 71.3 (2005), pp. 291–307.
- [98] Aleksandar Nikolov, Kunal Talwar, and Li Zhang. “The geometry of differential privacy: the sparse and approximate cases”. In: *Proceedings of the forty-fifth annual ACM symposium on Theory of computing*. 2013, pp. 351–360.
- [99] Elizabeth Meckes. “Projections of probability distributions: A measure-theoretic Dvoretzky theorem”. In: *Geometric aspects of functional analysis*. Springer, 2012, pp. 317–326.
- [100] Jim Young, Patrick Graham, and Richard Penny. “Using Bayesian networks to create synthetic data”. In: *Journal of Official Statistics* 25.4 (2009), p. 549.
- [101] Ian Covert, Scott M Lundberg, and Su-In Lee. “Understanding global feature contributions with additive importance measures”. In: *Advances in Neural Information Processing Systems* 33 (2020), pp. 17212–17223.
- [102] Nicolas Papernot, Martín Abadi, Úlfar Erlingsson, Ian Goodfellow, and Kunal Talwar. “Semi-supervised Knowledge Transfer for Deep Learning from Private Training Data”. In: *Proceedings of the International Conference on Learning Representations*. 2017. URL: <https://arxiv.org/abs/1610.05755>.
- [103] Martin Benning and Martin Burger. “Modern regularization methods for inverse problems”. In: *Acta Numerica* 27 (2018), pp. 1–111.
- [104] Peter L Bartlett, Andrea Montanari, and Alexander Rakhlin. “Deep learning: a statistical viewpoint”. In: *Acta numerica* 30 (2021), pp. 87–201.
- [105] Alon Halevy, Peter Norvig, and Fernando Pereira. “The unreasonable effectiveness of data”. In: *IEEE Intelligent Systems* 24.2 (2009), pp. 8–12.
- [106] Pedro Domingos. “A few useful things to know about machine learning”. In: *Communications of the ACM* 55.10 (2012), pp. 78–87.
- [107] Shaeke Salman and Xiuwen Liu. “Overfitting mechanism and avoidance in deep neural networks”. In: *arXiv preprint arXiv:1901.06566* (2019).
- [108] Zeke Xie, Fengxiang He, Shaopeng Fu, Issei Sato, Dacheng Tao, and Masashi Sugiyama. “Artificial neural variability for deep learning: On overfitting, noise memorization, and catastrophic forgetting”. In: *Neural computation* 33.8 (2021), pp. 2163–2192.
- [109] Claudio Filipi Gonçalves dos Santos and João Paulo Papa. “Avoiding Overfitting: A Survey on Regularization Methods for Convolutional Neural Networks”. In: *ACM Computing Surveys (CSUR)* (2022).
- [110] David A Van Dyk and Xiao-Li Meng. “The art of data augmentation”. In: *Journal of Computational and Graphical Statistics* 10.1 (2001), pp. 1–50.
- [111] Sebastien C Wong, Adam Gatt, Victor Stamatescu, and Mark D McDonnell. “Understanding data augmentation for classification: when to warp?” In: *2016 international conference on digital image computing: techniques and applications (DICTA)*. IEEE. 2016, pp. 1–6.
- [112] Sima Behpour, Kris M Kitani, and Brian D Ziebart. “Ada: Adversarial data augmentation for object detection”. In: *2019 IEEE Winter Conference on Applications of Computer Vision (WACV)*. IEEE. 2019, pp. 1243–1252.
- [113] David A Van Dyk and Xiao-Li Meng. “The art of data augmentation”. In: *Journal of Computational and Graphical Statistics* 10.1 (2001), pp. 1–50.

- 908 [114] Georgios Douzas, Fernando Bacao, Joao Fonseca, and Manvel Khudinyan. “Imbalanced learning
909 in land cover classification: Improving minority classes’ prediction accuracy using the geometric
910 SMOTE algorithm”. In: *Remote Sensing* 11.24 (2019), p. 3040.
- 911 [115] Wei Feng, Wenjiang Huang, and Wenxing Bao. “Imbalanced hyperspectral image classification with
912 an adaptive ensemble method based on SMOTE and rotation forest with differentiated sampling
913 rates”. In: *IEEE Geoscience and Remote Sensing Letters* 16.12 (2019), pp. 1879–1883.
- 914 [116] Roweida Mohammed, Jumanah Rawashdeh, and Malak Abdullah. “Machine learning with over-
915 sampling and undersampling techniques: overview study and experimental results”. In: *2020 11th
916 international conference on information and communication systems (ICICS)*. IEEE. 2020, pp. 243–
917 248.
- 918 [117] Julio Hernandez, Jesús Ariel Carrasco-Ochoa, and José Francisco Martínez-Trinidad. “An empirical
919 study of oversampling and undersampling for instance selection methods on imbalance datasets”.
920 In: *Iberoamerican Congress on Pattern Recognition*. Springer. 2013, pp. 262–269.
- 921 [118] Bartosz Krawczyk. “Learning from imbalanced data: open challenges and future directions”. In:
922 *Progress in Artificial Intelligence* 5.4 (2016), pp. 221–232.
- 923 [119] Teuvo Kohonen. “Emergence of invariant-feature detectors in the adaptive-subspace self-organizing
924 map”. In: *Biological cybernetics* 75.4 (1996), pp. 281–291.
- 925 [120] Jing Zhou, Yanan Zheng, Jie Tang, Jian Li, and Zhilin Yang. “Flipda: Effective and robust data
926 augmentation for few-shot learning”. In: *arXiv preprint arXiv:2108.06332* (2021).
- 927 [121] Yaqing Wang, Quanming Yao, James T Kwok, and Lionel M Ni. “Generalizing from a few examples:
928 A survey on few-shot learning”. In: *ACM computing surveys (csur)* 53.3 (2020), pp. 1–34.
- 929 [122] Hari Prasanna Das, Ryan Tran, Japjot Singh, Xiangyu Yue, Geoffrey Tison, Alberto Sangiovanni-
930 Vincentelli, and Costas J Spanos. “Conditional synthetic data generation for robust machine
931 learning applications with limited pandemic data”. In: *Proceedings of the AAAI Conference on
932 Artificial Intelligence*. Vol. 36. 11. 2022, pp. 11792–11800.
- 933 [123] Talip Ucar, Ehsan Hajiramezanali, and Lindsay Edwards. “Subtab: Subsetting features of tabular
934 data for self-supervised representation learning”. In: *Advances in Neural Information Processing
935 Systems* 34 (2021), pp. 18853–18865.
- 936 [124] Tsubasa Takahashi, Shun Takagi, Hajime Ono, and Tatsuya Komatsu. “Differentially Private Vari-
937 ational Autoencoders with Term-wise Gradient Aggregation”. In: *arXiv preprint arXiv:2006.11204*
938 (2020).
- 939 [125] Katherina K Hauner, Richard E Zinbarg, and William Revelle. “A latent variable model approach
940 to estimating systematic bias in the oversampling method”. In: *Behavior Research Methods* 46.3
941 (2014), pp. 786–797.